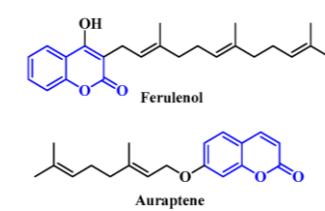
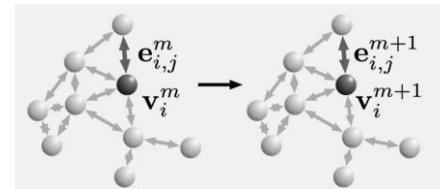
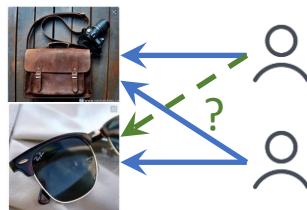


# Graph Neural Networks Applications

Rex Ying



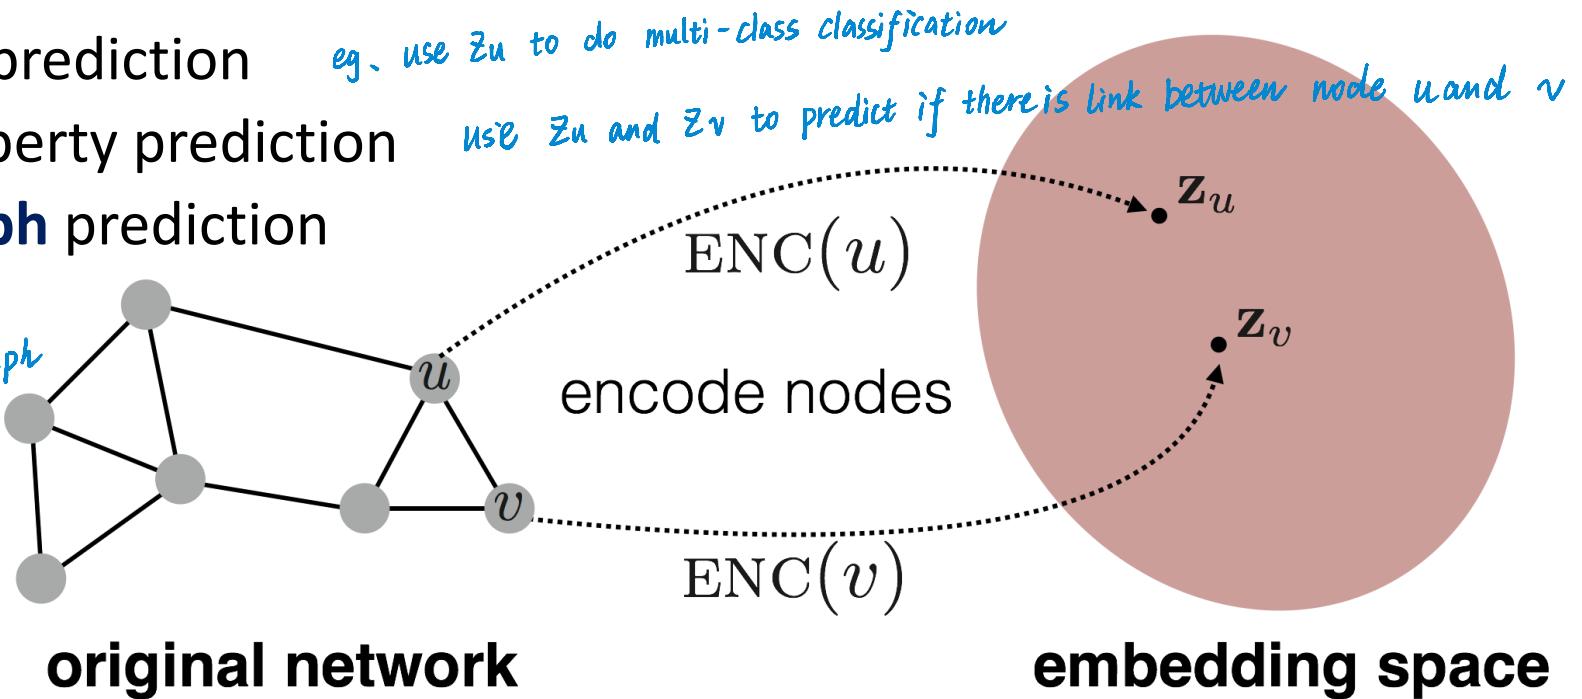
# Background: Graph Representation Learning

- Given one or more input graphs, we use a **(deep) encoder** to map nodes to **high-dimensional** embedding space  $u \rightarrow z_u$

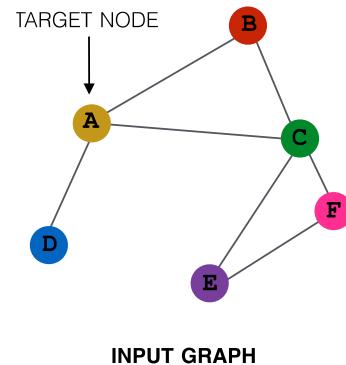
- Objectives

- Node** property prediction
- Link / edge** property prediction
- Graph / subgraph** prediction

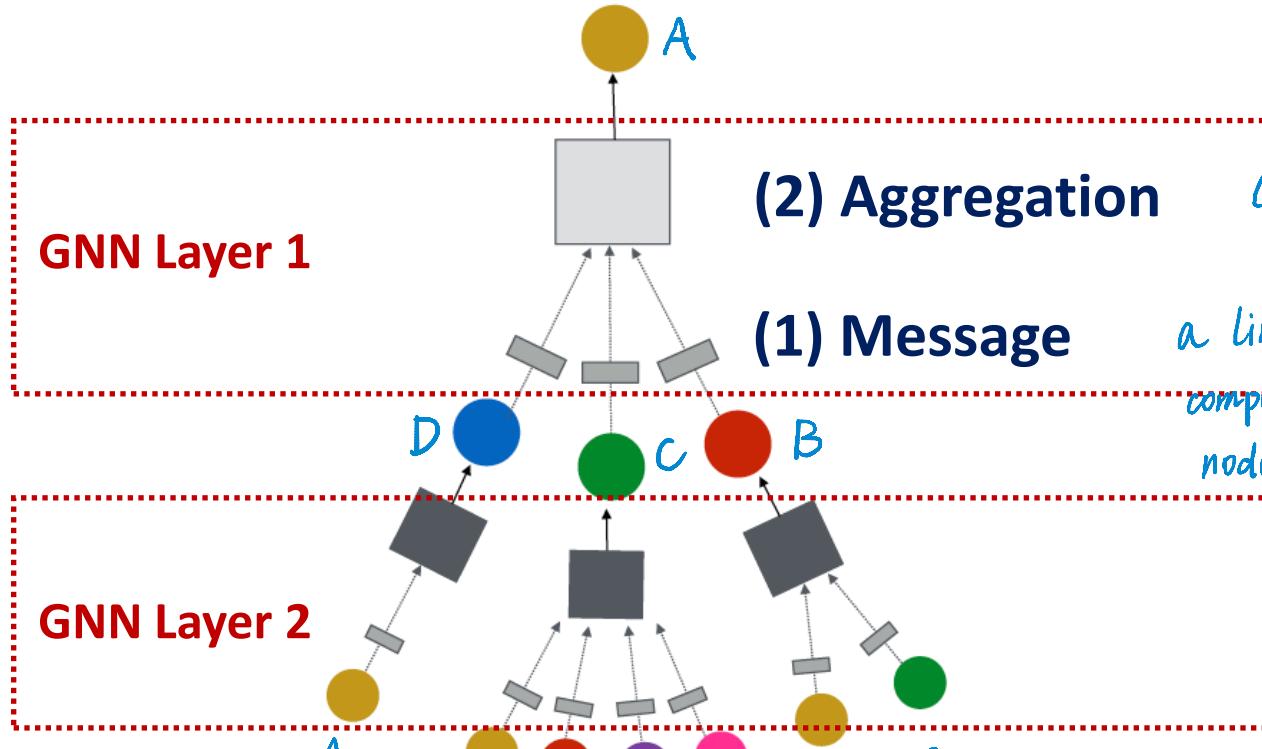
average over all node  
embeddings to one graph  
embedding  $\rightarrow$  make  
prediction



# Recap: Graph Neural Networks (GNNs)



**(3) Layer connectivity**

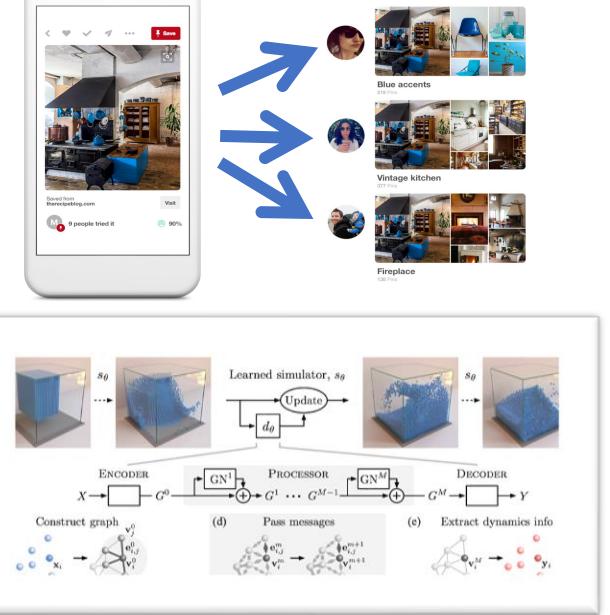


# GNN Applications

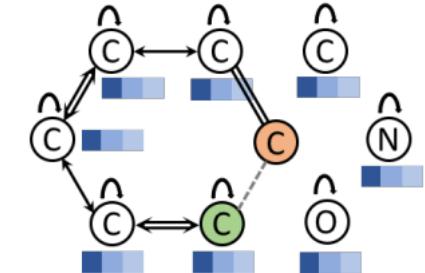
- Social Networks
- Natural Science
- Medicine
- Explainability of GNNs



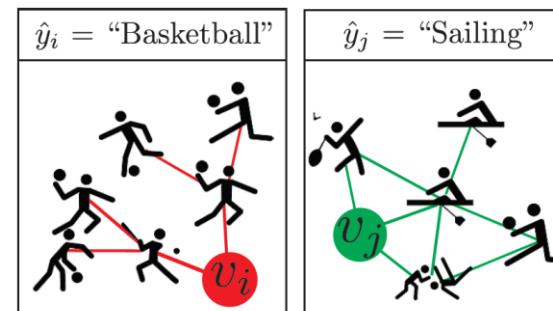
Recommender System



Physical Simulation

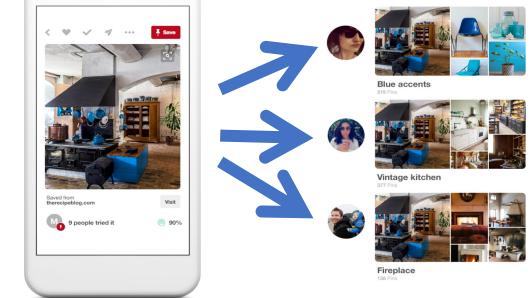


Drug Side-effects  
Molecule Generation



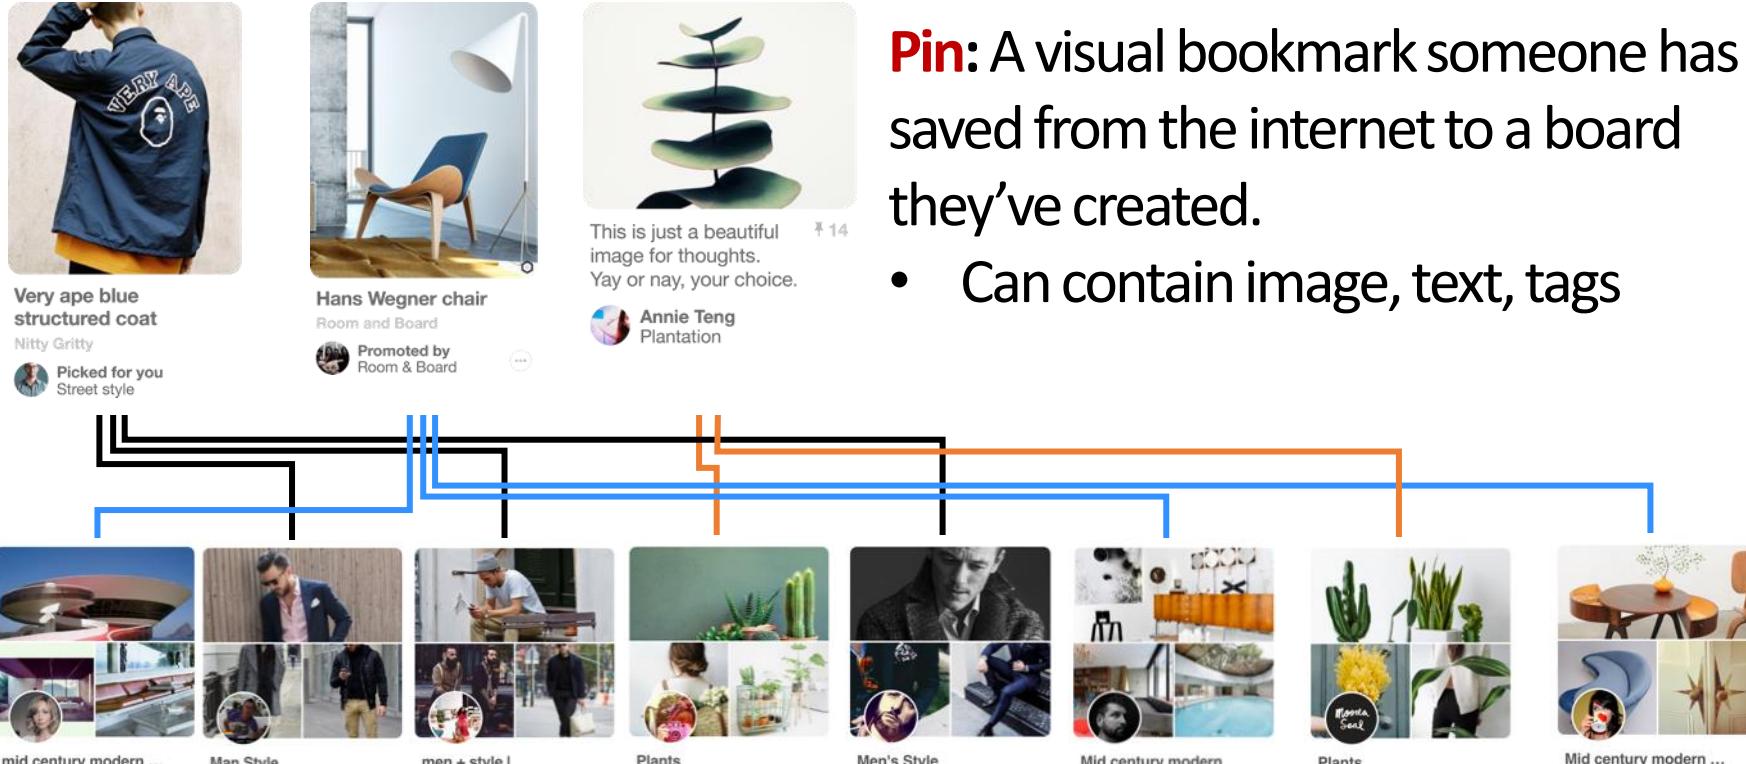
# GNN Applications

- Social Networks → Recommender System
- Natural Science
- Medicine
- Explainability of GNNs



# Pinterest Recommender System

## Human curated collection of pins



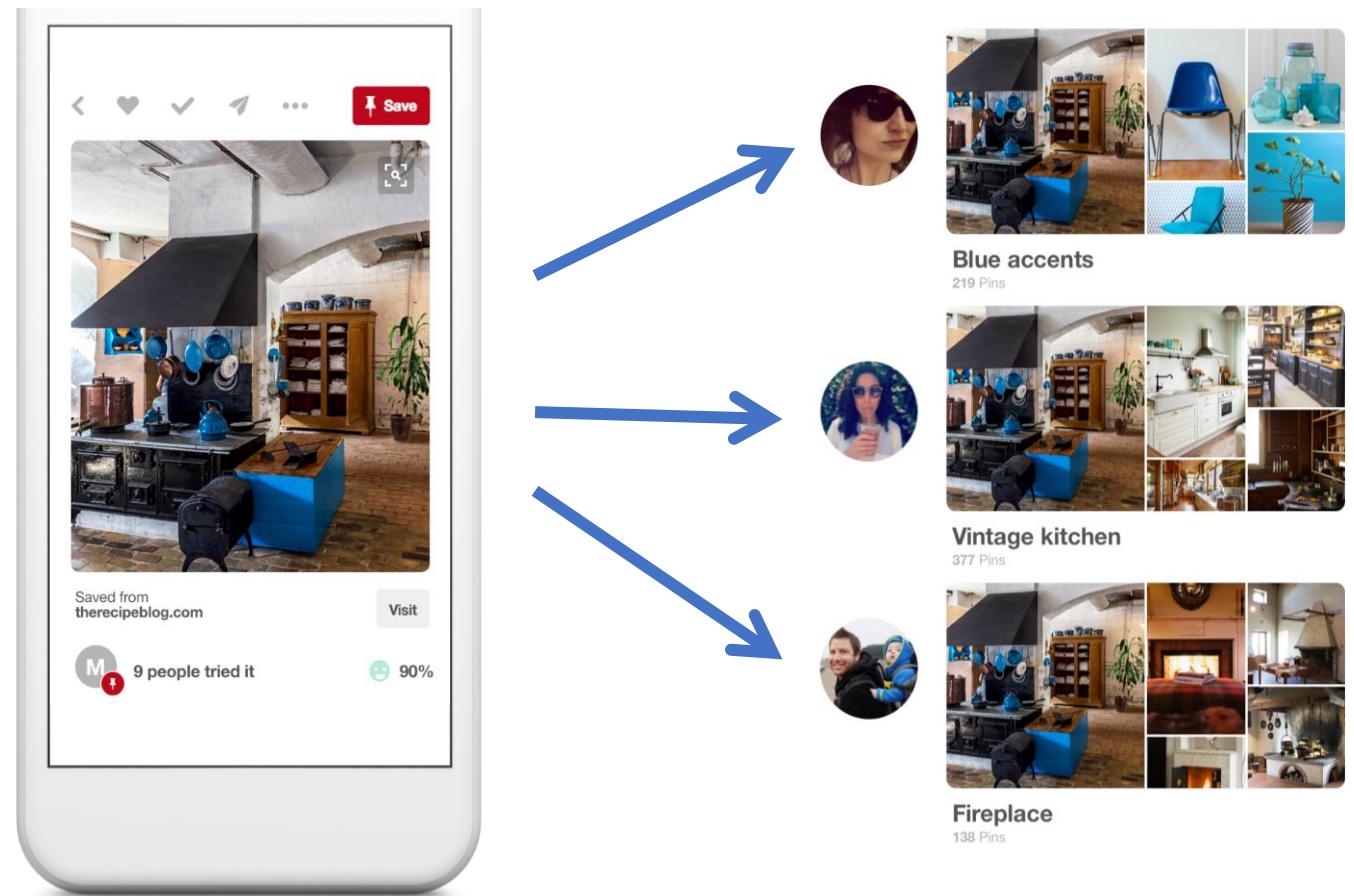
**Pin:** A visual bookmark someone has saved from the internet to a board they've created.

- Can contain image, text, tags

**Board:** A collection of ideas (pins having something in common)

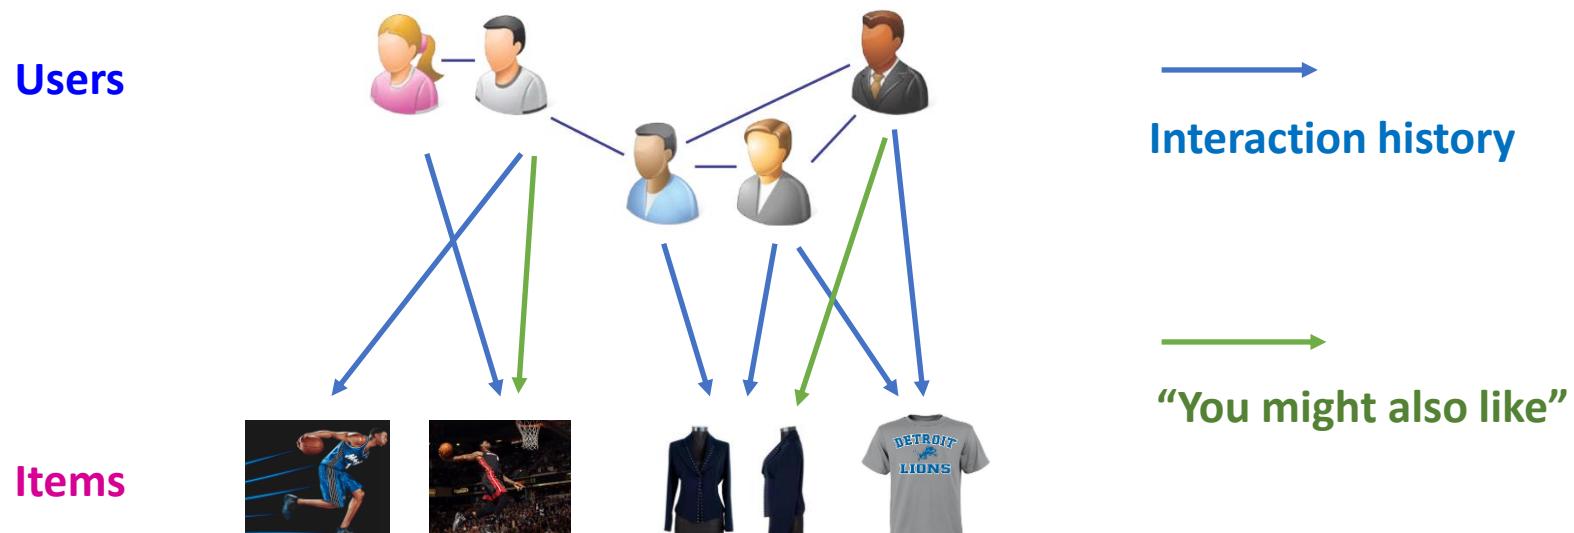
# Pinterest Recommender System

- Social network for idea collections
- 300M users
- 4B+ pins, 2B+ boards

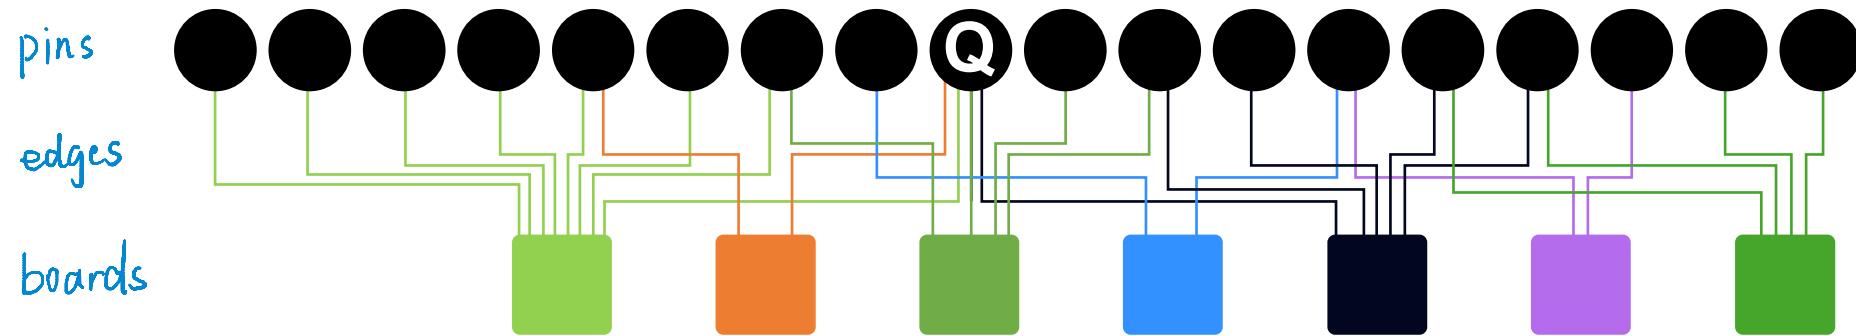


# PinSage: Web-scale Recommender System

- Users are related by social network
- Users interact with items (watch movie, buy merchandise, listen to music)
- Predict future interactions from history



# Pinterest Graph

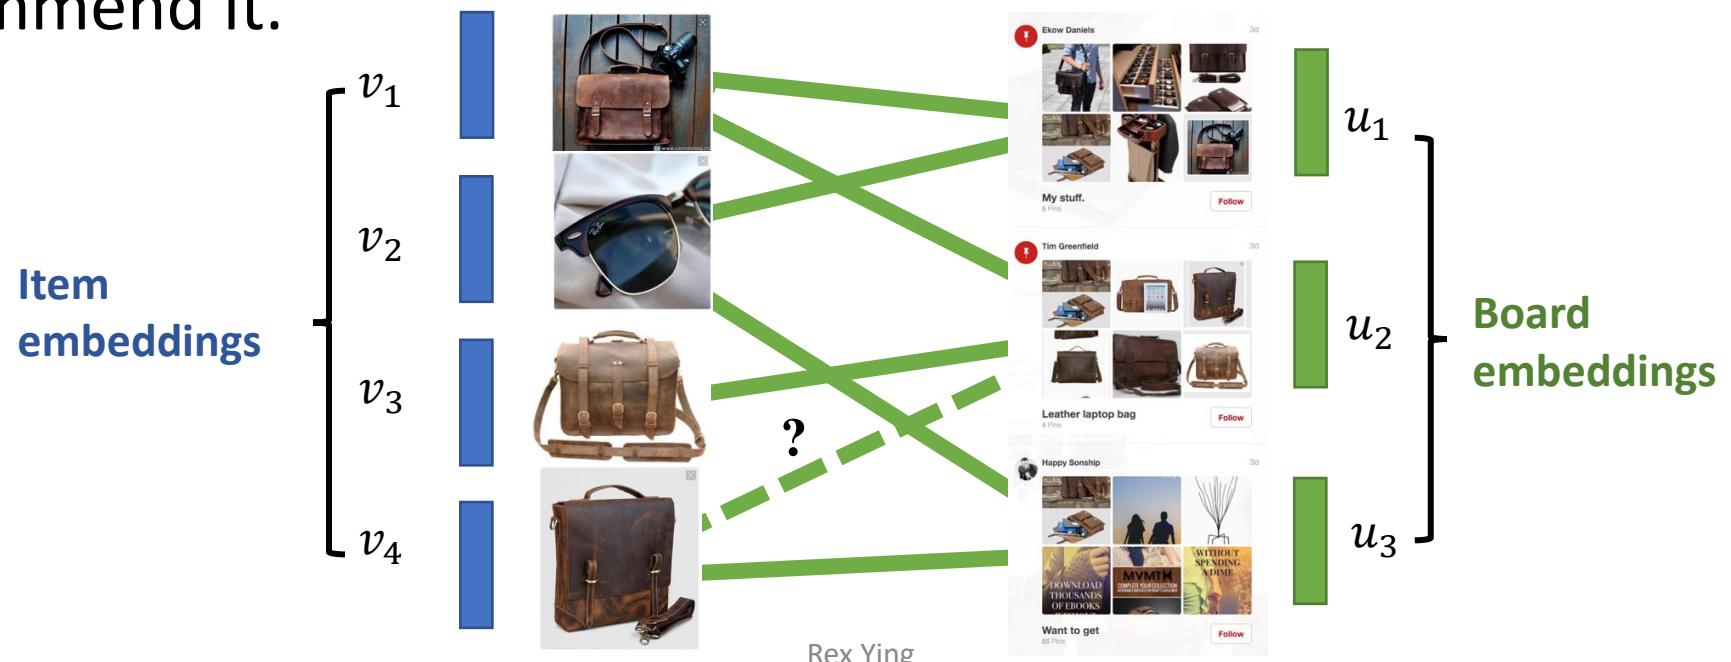


**Graph:** 2B pins, 1B boards, 20B edges

- **Graph is dynamic:** Need to apply to new nodes without model retraining
- **Rich node features:** Content, images

# Recommend by embeddings

- Learn embeddings for **items** and **boards** (*users*)
- **Query:** which item to recommend to the board with embedding  $u_2$  ?
- **Answer:** find the closest embedding ( $v_4$ ) by nearest neighbor. *in item embeddings*  
Recommend it.



# Why is it Hard? *Scalability*

How to scale the training as well as inference of node embeddings to  
**graphs with billions of nodes and tens of billions of edges?**

- Scaling up is difficult:
  - Existing collaborative filtering and distributed node embedding methods are inefficient when the underlying graph has billions of nodes and whose structure is constantly evolving

# PinSage Overview

- **PinSage** graph convolutional network:
  - **Goal:** Generate embeddings for nodes (e.g., Pins/images) in a web-scale Pinterest graph containing billions of objects
  - **Key Idea:** Borrow information from nearby nodes
    - E.g., bed rail Pin might look like a garden fence, but gates and beds are rarely adjacent in the graph
  - Pin embeddings are essential to various tasks like recommendation of Pins, classification, clustering, ranking
    - Services like “Related Pins”, “Search”, “Shopping”, “Ads”

gates pin and beds pin  
are in different boards  
though they look similar in CV  
so we need graph structure  
to model pins relationship  
and boards

# PinSage Pipeline

1. **Collect** billions of training pairs from logs.
  - **Positive pair**: a pin and a board that contains it
  - **Negative pair**: a **random** pair of pin and board
    - With high probability the pin is not pinned to the board
2. **Train GNN** to generate similar embeddings for training pairs.
3. **Inference**: generate embeddings for all pins.
4. **Nearest neighbor search** in embedding space to make recommendations.

# Neighborhood Sampling

## Constructing convolutions via random walks

- Performing convolutions on full neighborhoods is infeasible:  
*bc we can't fit the whole graph to GPU . we have to do mini-batch training , look at few nodes and its neighbors  
only pass selected neighbor nodes to GPU and compute node embeddings*
- How to select the set of neighbors of a node to convolve?
- **Personalized PageRank** can help!  
*what are the important nodes (pins) for a node (board)*
- **Importance pooling:** Define importance-based neighborhoods by simulating random walks and selecting the neighbors with the highest visit counts

# Training Objective

- Train so that **pins that are consecutively clicked have similar embeddings.**
- Max-margin loss:

$$\mathcal{L} = \sum_{(u,v) \in \mathcal{D}} \max(0, -\mathbf{z}_u^\top \mathbf{z}_v + \mathbf{z}_u^\top \mathbf{z}_n + \Delta)$$

ad: simple computation is efficient in training  
similarity is inner product of node embedding  $\mathbf{z}_u$  and  $\mathbf{z}_v$   
pos pair will differ from neg pair by margin

set of training pairs from user logs

“positive”/true training pair  $u$  and  $v$   $\approx 1$

“negative” example  $u$  and  $n$   $\approx -1$

“margin” (i.e., how much larger positive pair similarity should be compared to negative)  $0.5$

# Negative Sampling

- **Hard negative sampling**

刺绣 embroidery  
十字绣 cross stitch



Query



Positive Example

帽子 real



Random Negative

bird 刺绣



Hard Negative

same texture 刺绣  
but not same object (bird)

(confusing example)

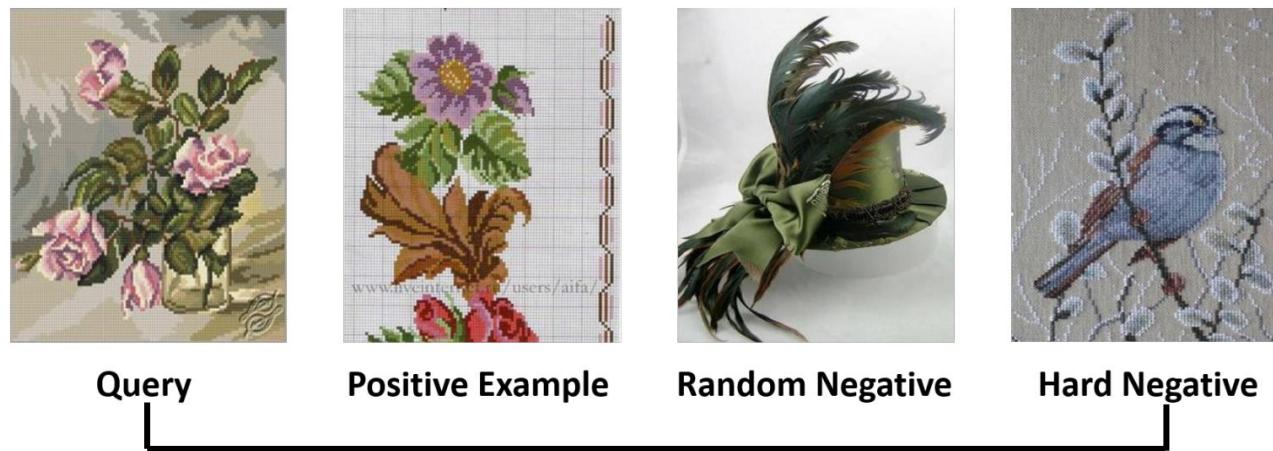
- Use personalized PageRank (PPR)

- Use nodes that have PPR score ranked at 1000-5000 as hard negatives
- Have something in common, but not too similar

Harder to distinguish from positive

# Negative Sampling

- **Hard negative sampling**



Harder to distinguish from positive

- **Curriculum training** on hard negatives

- provide harder and harder examples over time

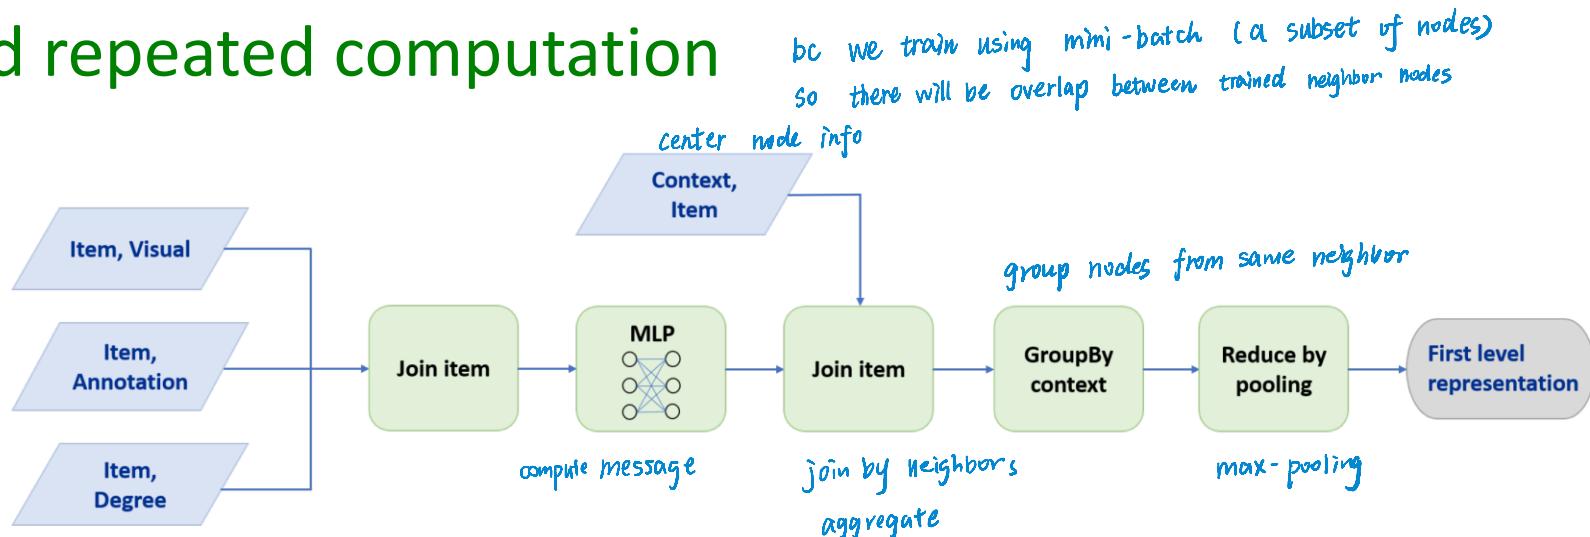
After training

# Efficient Inference

## Efficient inference via **MapReduce**

- Bottom-up aggregation of node embeddings lends itself to MapReduce
  - Decompose each aggregation step across all nodes into three operations in MapReduce, i.e., *map*, *join*, and *reduce*
- Avoid repeated computation

For every layer:



# Evaluation

- Baselines: *(pretrained)*

- **Visual**: VGG visual embeddings for recommendations
- **Annotation**: Word2vec embeddings
- **Combined**: Concatenate embeddings:
  - Uses exact same data and loss function as PinSage

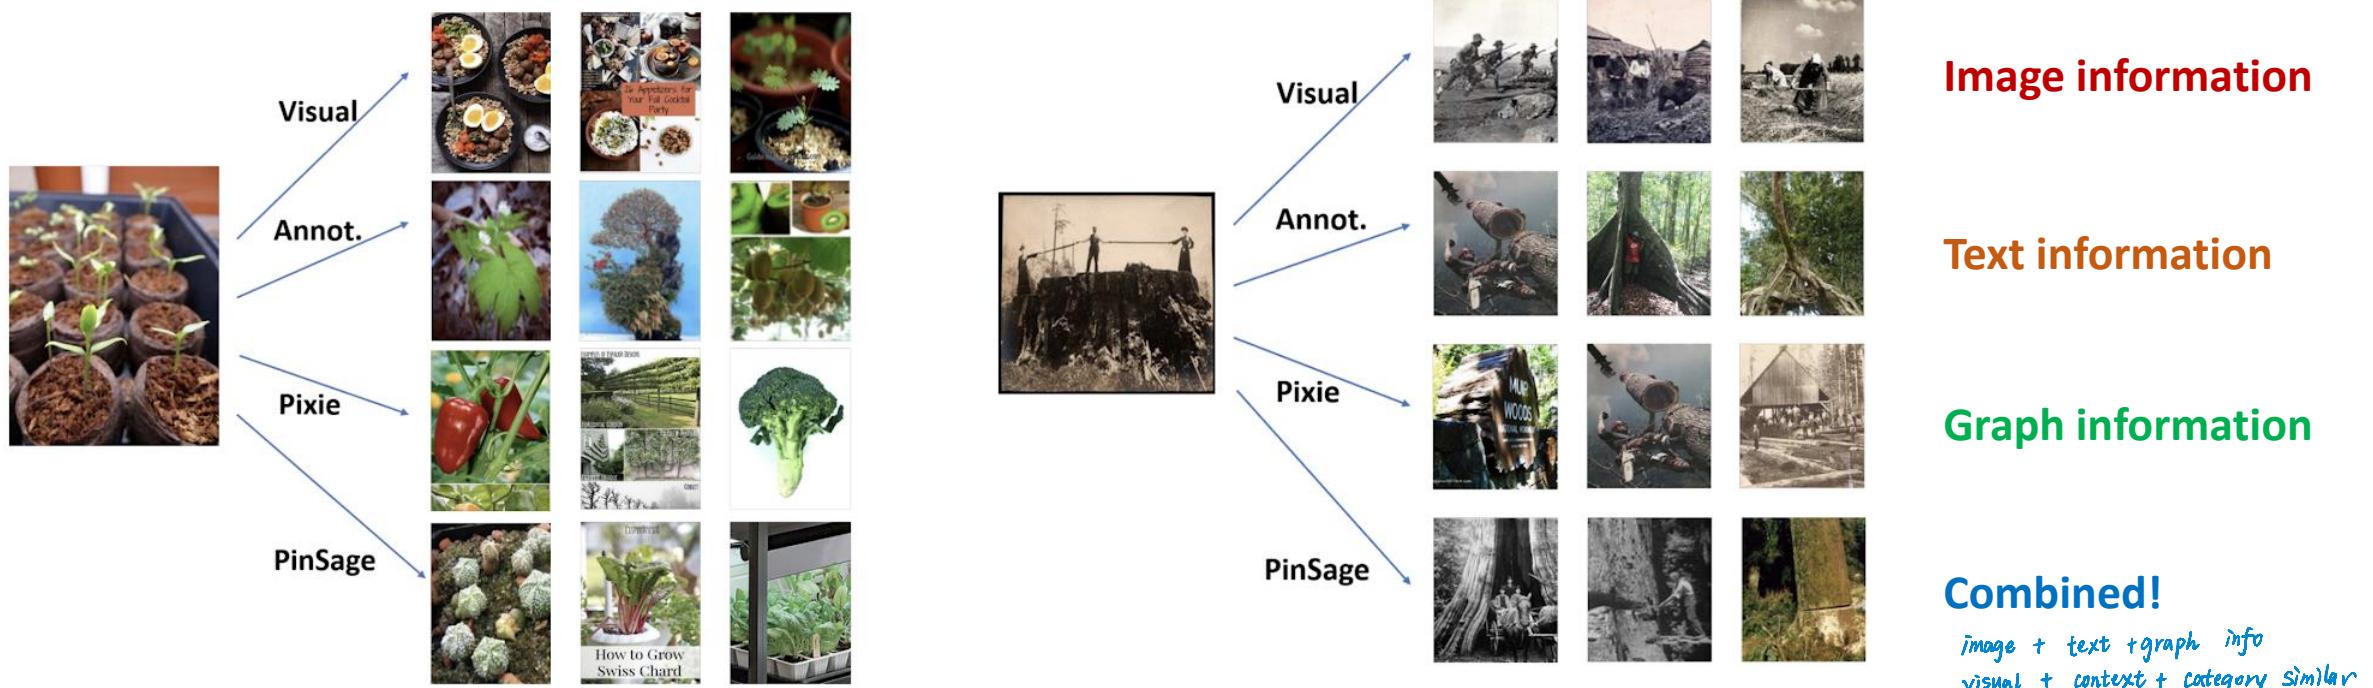


*recommend based on  
near neighbor*

Method	Hit-rate	MRR
Visual	17%	0.23
Annotation	14%	0.19
Combined	27%	0.37
max-pooling	39%	0.37
mean-pooling	41%	0.51
mean-pooling-xent	29%	0.35
mean-pooling-hard	46%	0.56
PinSage	67%	<b>0.59</b>

**PinSage gives 150% improvement in hit rate and 60% improvement in MRR over the best baseline**

# Evaluation



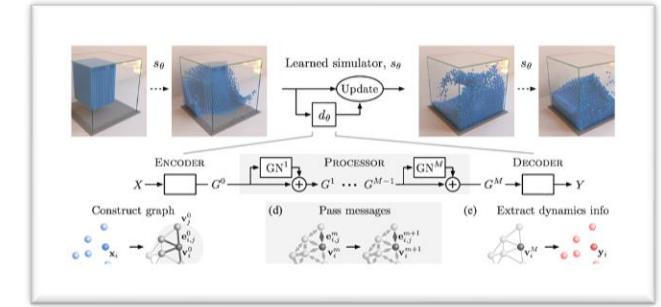
Pixie is a purely graph-based method that uses biased random walks to generate ranking scores by simulating random walks starting at query Pin. Items with top scores are retrieved as recommendations [Eksombatchai et al., 2018]

# Summary

- Many online platforms can be modeled as graphs, where nodes are users, items and edges represent interactions
  - View, purchase, review, follow, likes
- We use link prediction objective to train a GNN-based recommender systems
  - Efficient training and inference on a billion-scale graph
- The learned graph embedding captures image/text feature as well as graph structure (category feature)
- The learned embedding is useful in many downstream tasks beyond recommendation!

# GNN Applications

- Social Networks
- Natural Science  Physical Simulation
- Medicine
- Explainability of GNNs

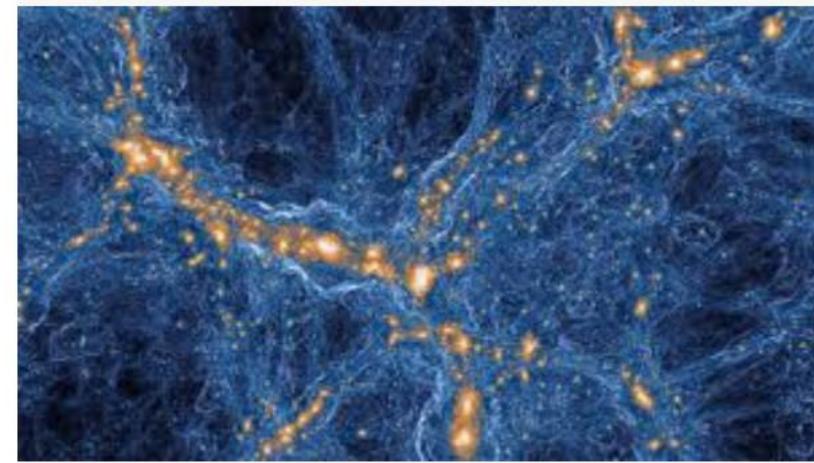


# Simulations in Science and Engineering

## 1. Particle-particle interactions:



Water simulation



Galaxy formation

node : stars  
edge : gravity force

We construct graphs with particles as nodes and interactions as edges

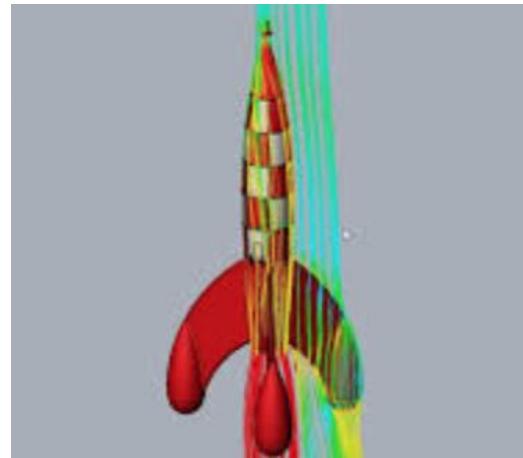
# Simulations in Science and Engineering

## 2. PDEs (on grid or mesh)

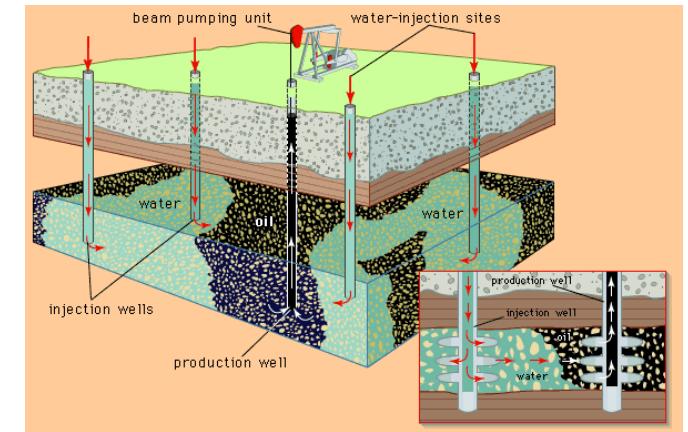
Partial Differentiation Equation



Weather prediction



Aerodynamics



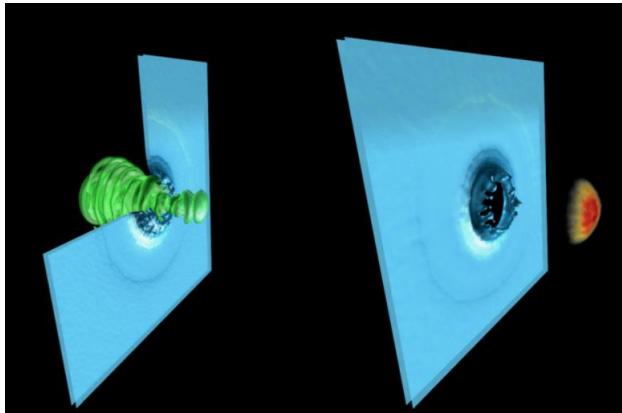
Reservoir simulation

Graph based on grid / mesh structure

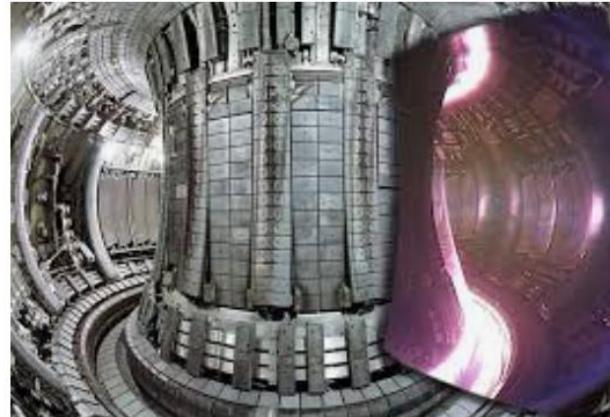
alternative 3-D U-net

# Simulations in Science and Engineering

## 3. Particle-in-Cell (involves both grid and particles)

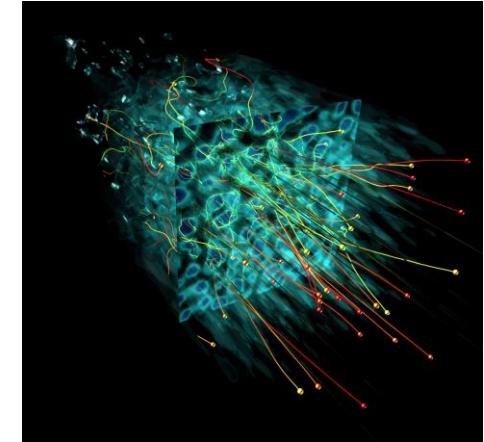


Laser-plasma particle  
acceleration  
激光等离子加速



Fusion  
核聚变

Graph combining grid / mesh with particles!



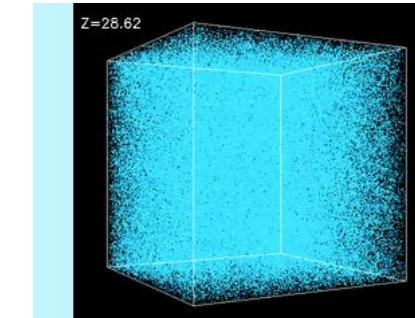
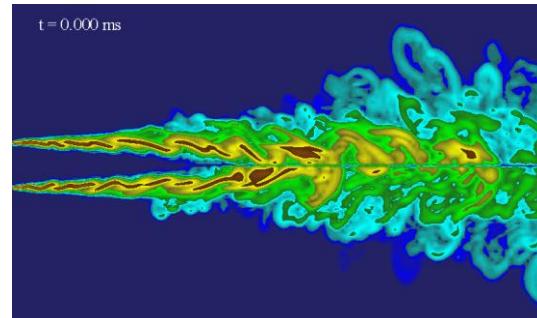
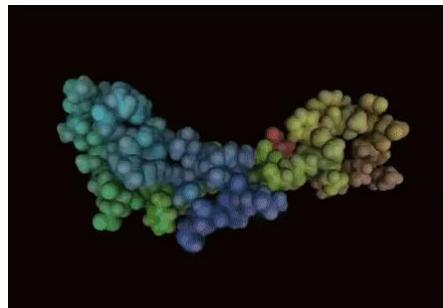
Cosmic-ray acceleration  
宇宙射线加速

# Why Learning Simulation

## Engineered simulators:

1. Substantial effort to build
2. Substantial resources to run
3. Only as accurate as the designer
4. Not always suitable for solving inverse problems

*eg inverse problem: How to tune params to stabilize particles for 10s?*



## Learned simulators:

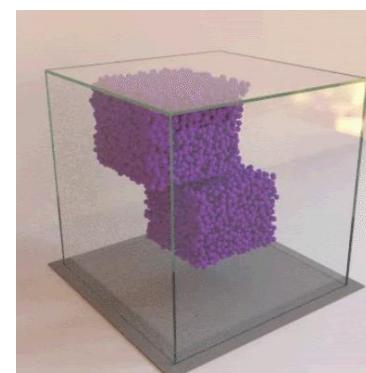
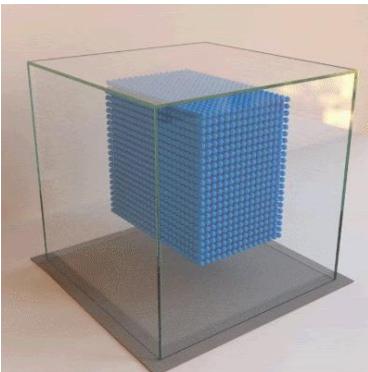
1. Shared architectures *bc generalizability*
2. Can be directly optimized for **efficiency**
3. Can be as accurate as the available **data**
4. Gradient-based search for **control**

*use gradient to backtrace param*

# Dynamic predictions – Learning Simulation

- Simulating complex fluids and other materials:

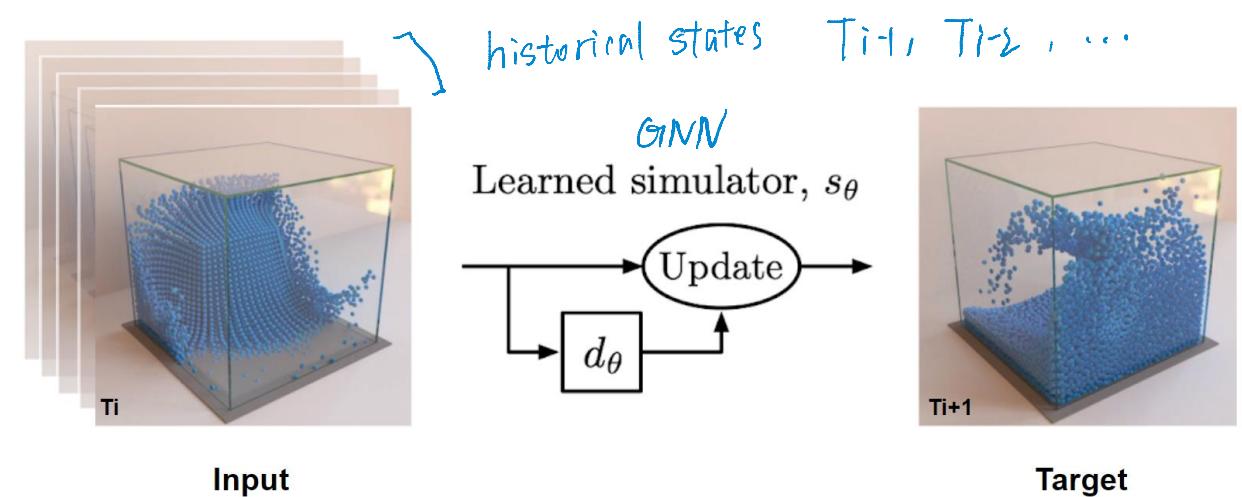
initial state  $\rightarrow$  future state



# Use historical particle information

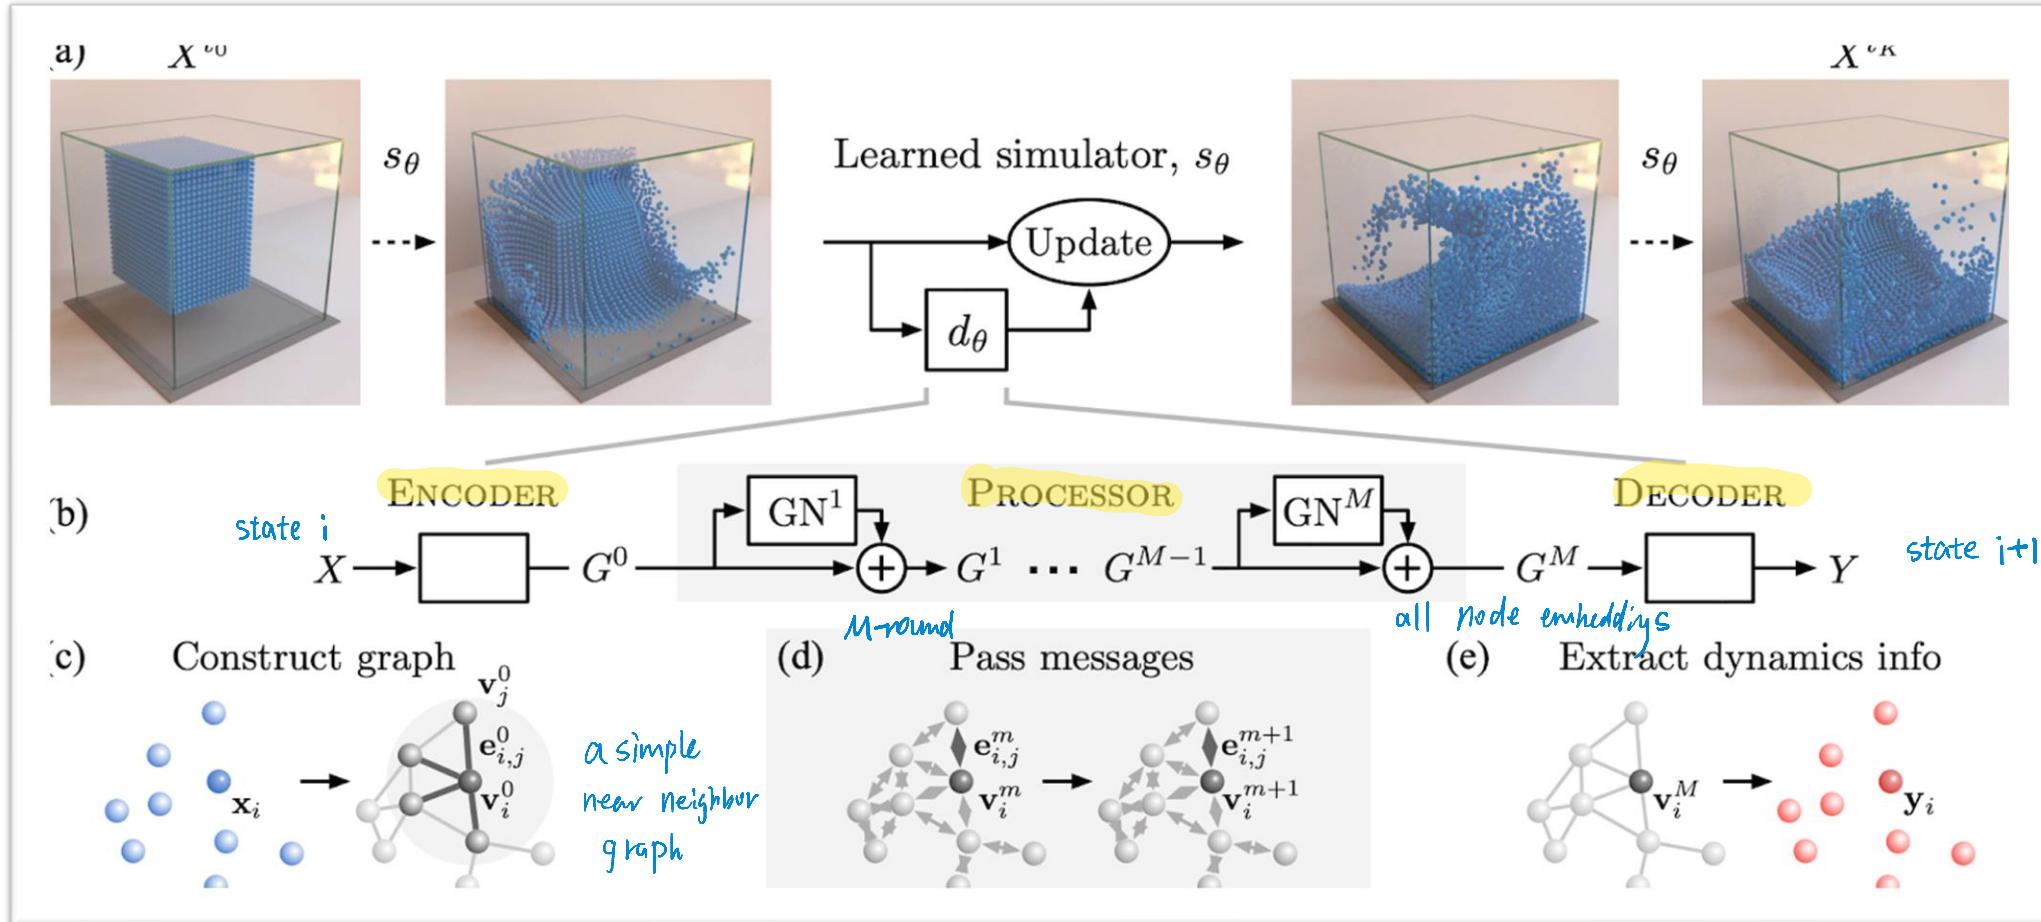
- During training, we are given **particle properties** (position, velocity ...) for a time period  $0 - T$ 
  - Equal and discrete time interval
- We **sample** batches of particle states at consecutive time steps
- Historical particle information

**Multiple instances per minibatch**



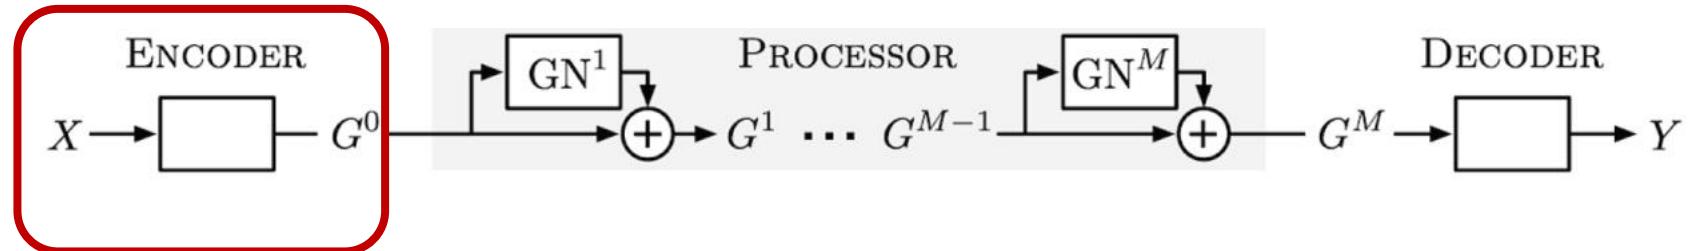
# Simulation Architecture

{ training : state  $t_i$  and  $t_{i+1}$  pairs  
 inference : input state  $s_\theta$ , can predict state at any time point



# Graph Network Simulator (GNS) Model

## Encoder

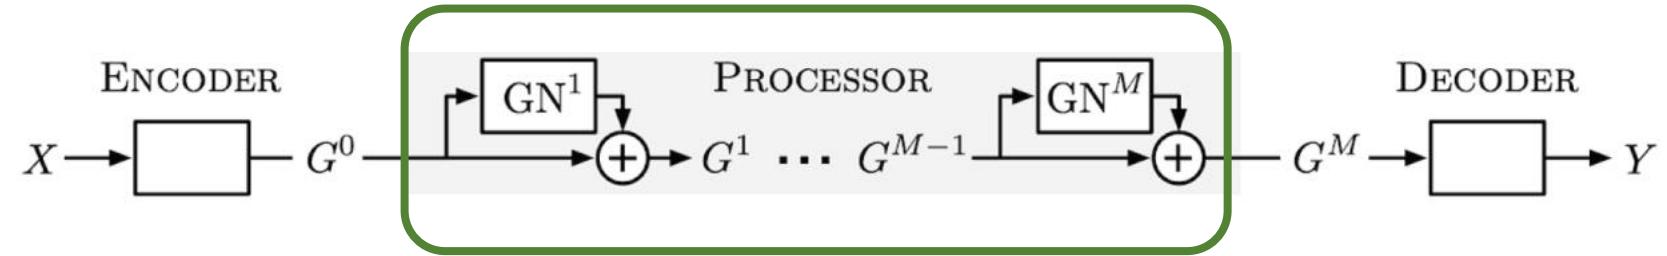


- Node input features:
  - Position, previous 5 velocities, particle type (water/sand)
- Edge input features: displacements
  - Embed features with MLP
- Construct neighbourhood graph
  - K-nearest neighbor (kNN) edges
    - choice depends on inductive bias for specific problem

Learning to Simulate Complex Physics with Graph Networks, ICML 2020

# Graph Network Simulator (GNS) Model

- **Processor**

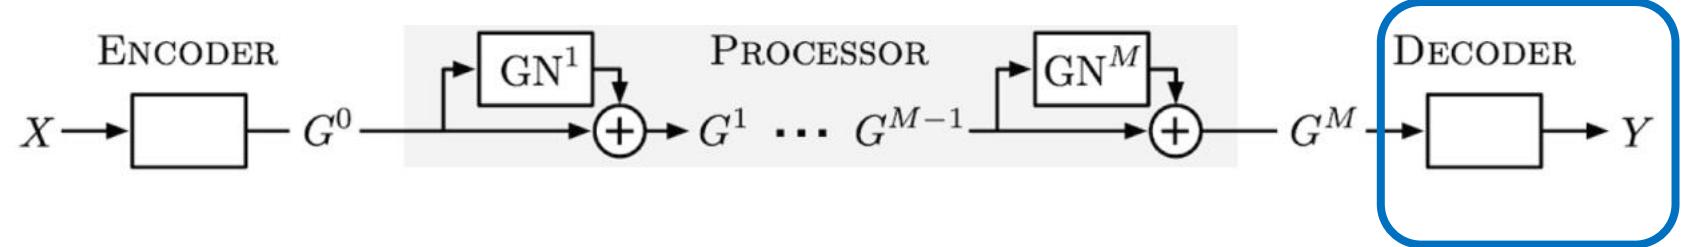


- Message-passing layers (x10) on the kNN graph

- Concatenate node and edge features, and compute message function through multi-layer perceptron (MLP)
- Outputs embeddings for each particle
  - Used to predict next step dynamics

# Graph Network Simulator (GNS) Model

**Decoder**

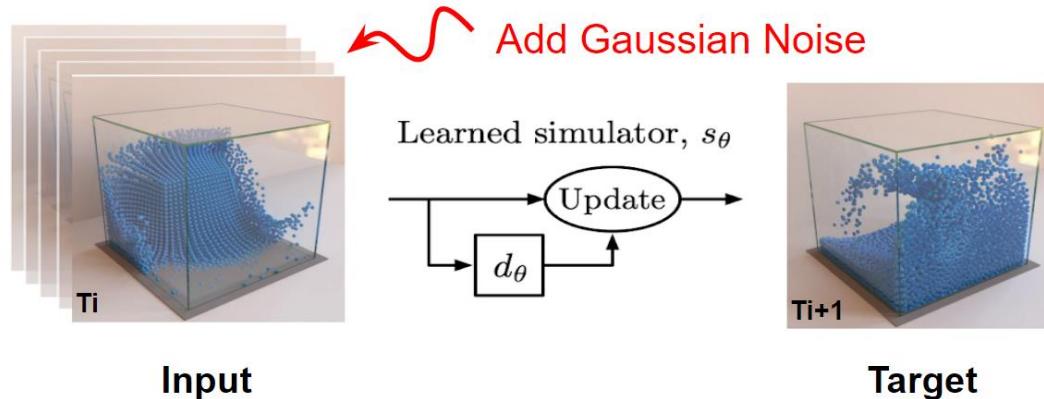


- Decode **acceleration** 加入速度  $v$
- Feed into Euler integrator to obtain position and velocity
- Sum over L2 Loss for all particles in each pair

# Graph Network Simulator (GNS) Model

- Training time: One-step minibatch training

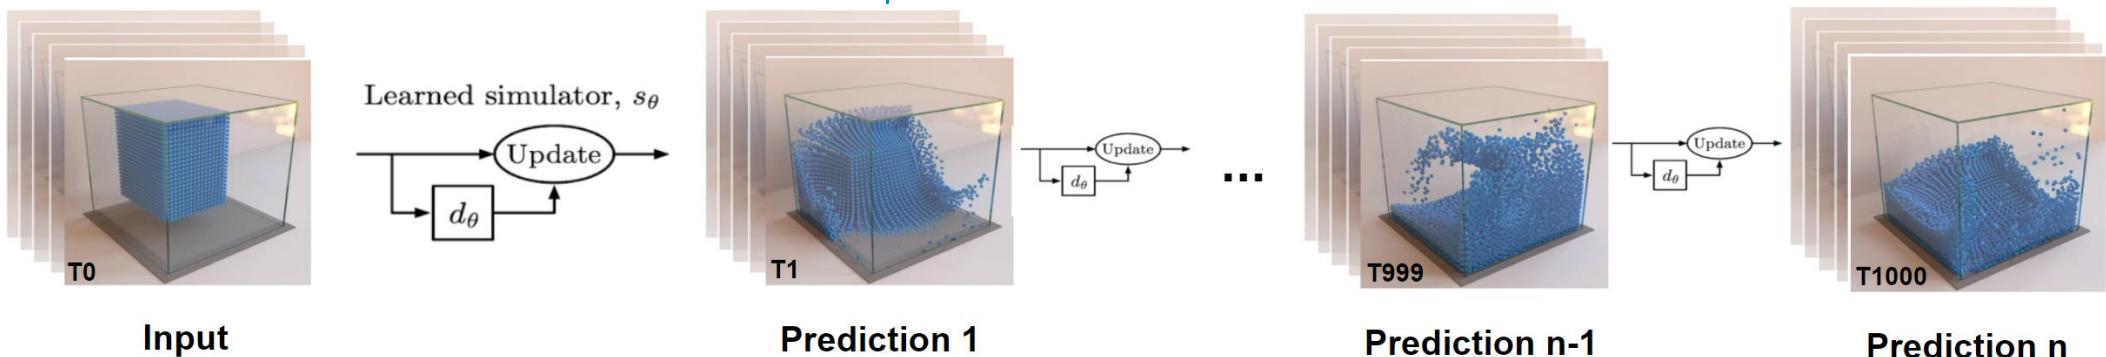
$t_i - t_{i+1}$  pair



Gaussian  
Use noise at training to prevent  
error accumulation at evaluation  
make GNN robust  
variant: use Recurrent GNN

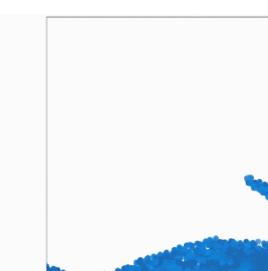
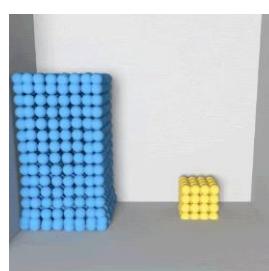
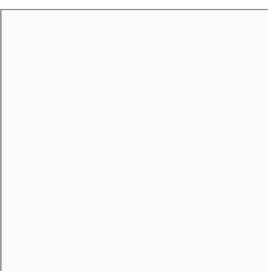
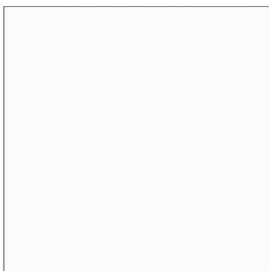
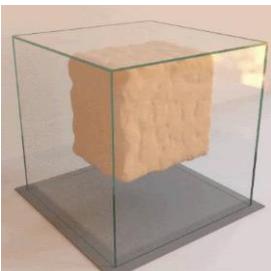
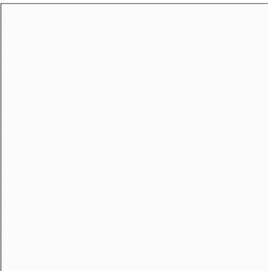
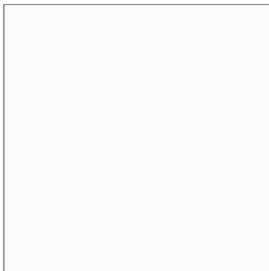
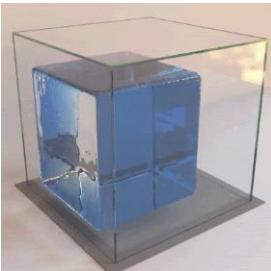
- Test time: 1000s of steps

prediction serve as input



# Generalization

- Same model and hyperparameters across datasets



local interaction is key to  
simulation

# Summary

- In simulation, we are given an initial condition, and use the model to make predictions of the **evolution of the system** over time.
- We model the systems through **particles** (nodes) and **interactions** between particles (edges).
- The model **generalizes** to many different scenarios since the model learns the underlying rules of physical interactions.
- Open question
  - more complex interactions
  - more efficient simulation of large systems

# GNN Applications

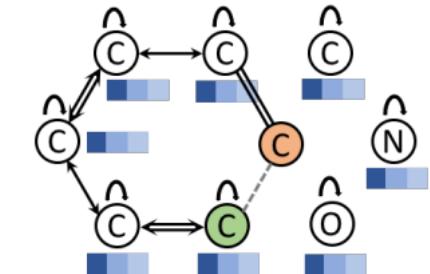
- Social Networks

- Natural Science

- Medicine

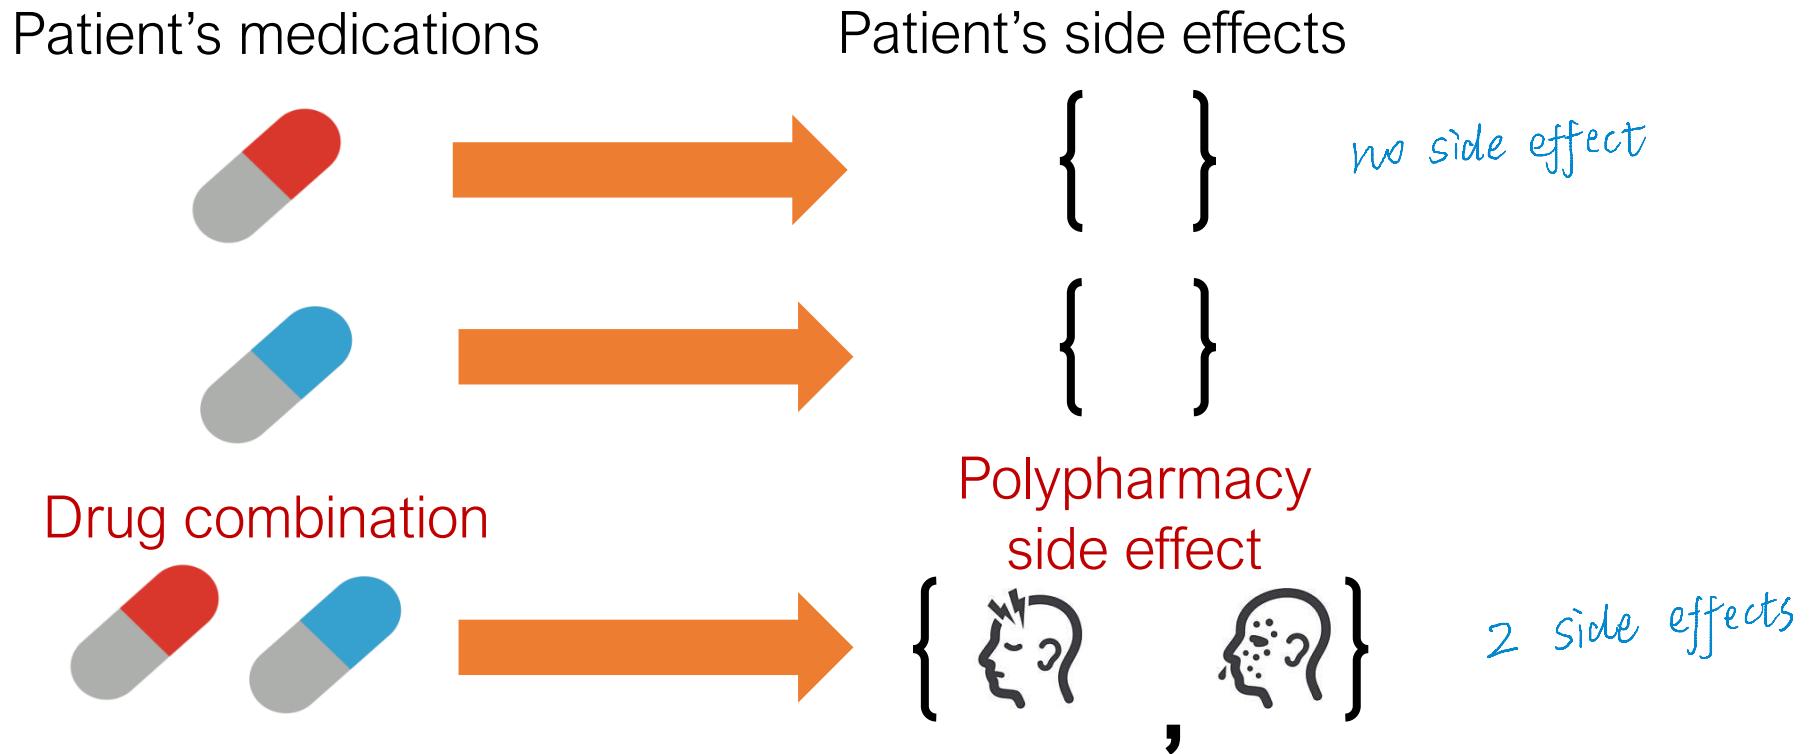


**Drug Side-effects  
Molecule Generation**



- Explainability of GNNs

# Polypharmacy Side Effects



Polypharmacy: use multiple drugs for a disease

Modeling Polypharmacy Side Effects with Graph Convolutional Networks, Bioinformatics 2018

# Polypharmacy Side Effects

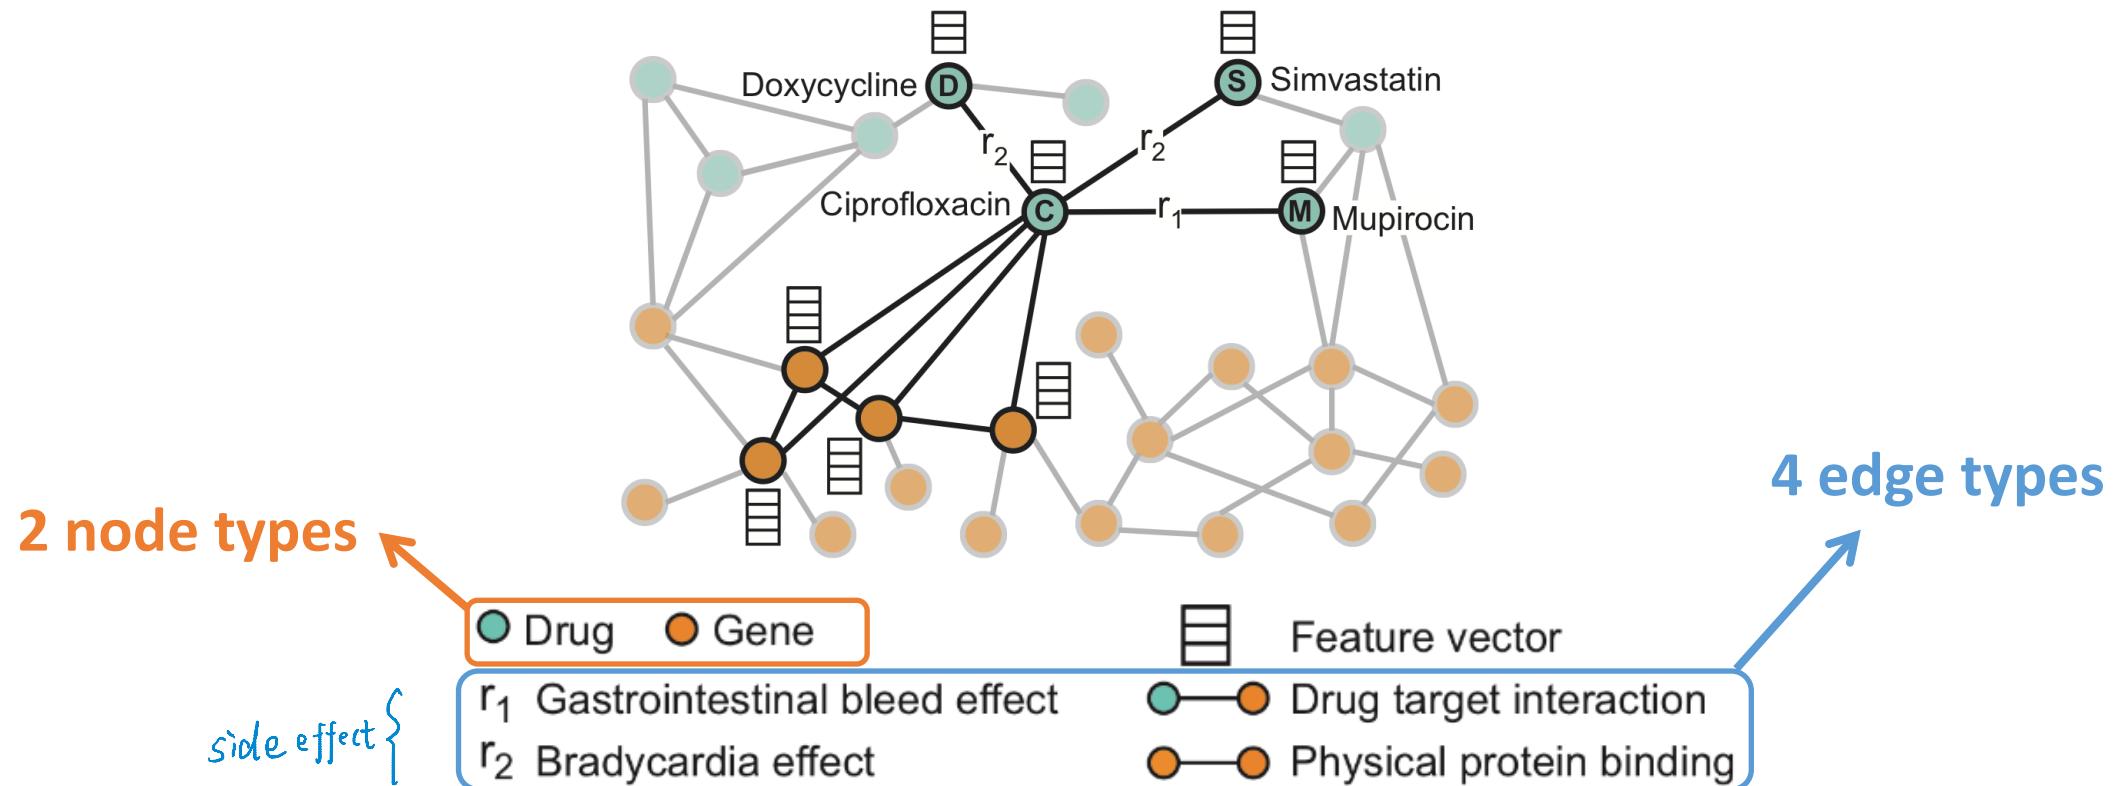
- Polypharmacy is common to treat complex diseases and co-existing conditions
- High risk of side effects due to interactions
- 15% of the U.S. population affected
- Annual costs exceed \$177 billion
- Difficult to identify manually:
  - Rare, occur only in a subset of patients
  - Not observed in clinical testing

# Modeling Polypharmacy

- Systematic experimental screening of drug interactions is challenging
- Idea: Computationally screen/predict polypharmacy side effects
  - Use molecular, pharmacological and patient population data
  - Guide translational strategies for combination treatments in patients

# Data: Heterogeneous Graphs

- **Heterogeneous (multimodal) graphs:** graphs with different node types and/or edge types



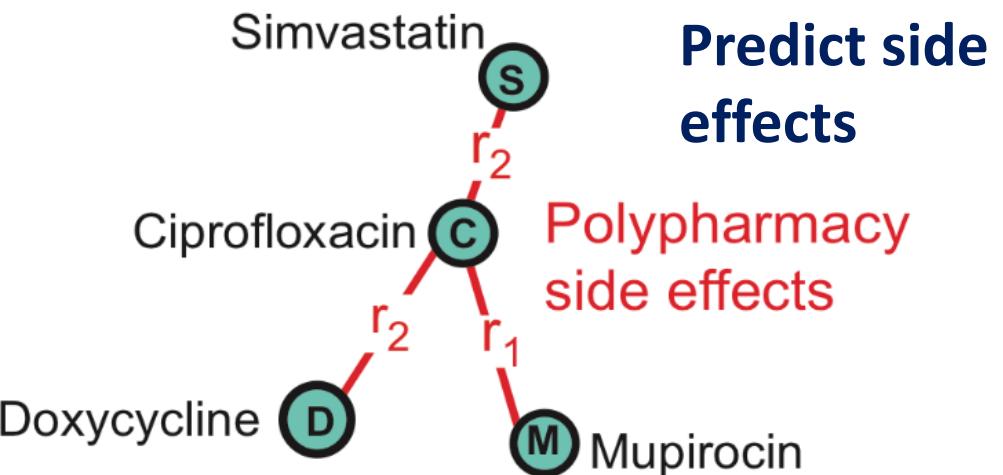
# Task Description

## Link Prediction

- Predict **labeled edges** between drugs nodes
  - i.e., predict the likelihood that an edge  $(c, r_2, s)$  exists between drug nodes  $c$  and  $s$
  - **Meaning:** Drug combination  $(c, s)$  leads to polypharmacy side effect  $r_2$

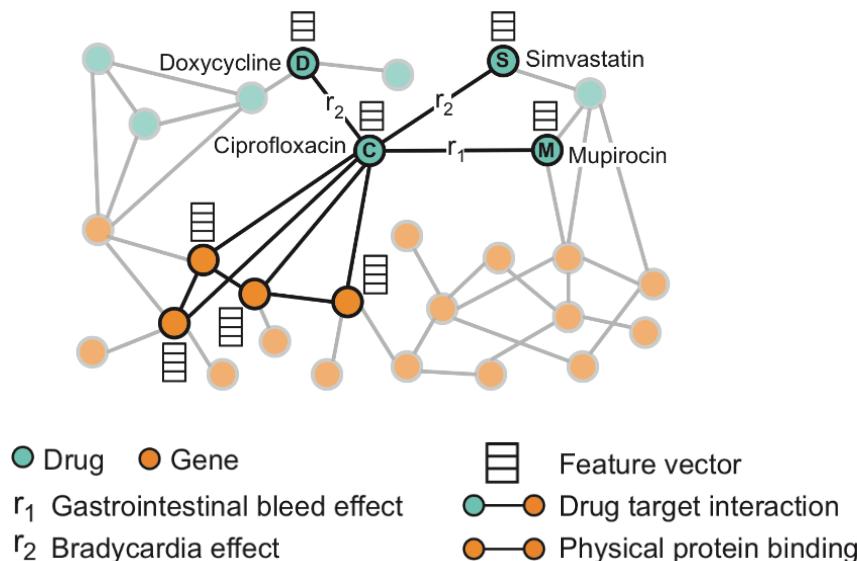
● Drug   ● Gene  
 $r_1$  Gastrointestinal bleed effect  
 $r_2$  Bradycardia effect

■ Feature vector  
— Drug target interaction  
— Physical protein binding



# Model: Heterogenous GNN

- **Key Insight:** Compute GNN messages from each edge type, then aggregate across different edge types
  - **Input:** heterogenous graph
  - **Output:** node embeddings



## One layer of Heterogeneous GNN

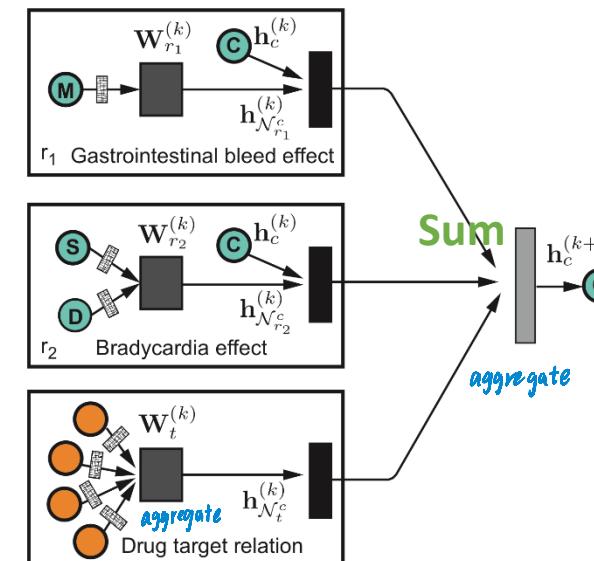
GNN for  
Edge type:

$r_1$

GNN for  
Edge type:  
 $r_2$

$r_2$

GNN for  
Edge type:  
drug-target



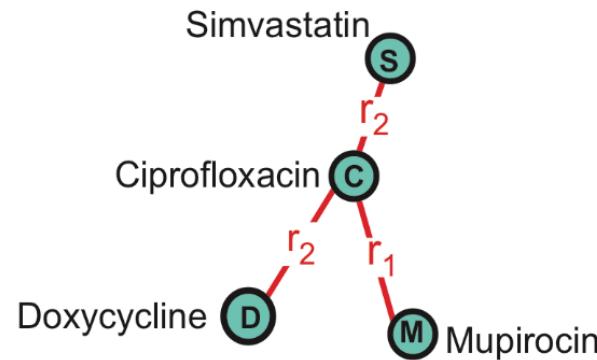
node embedding

# Making Edge Predictions

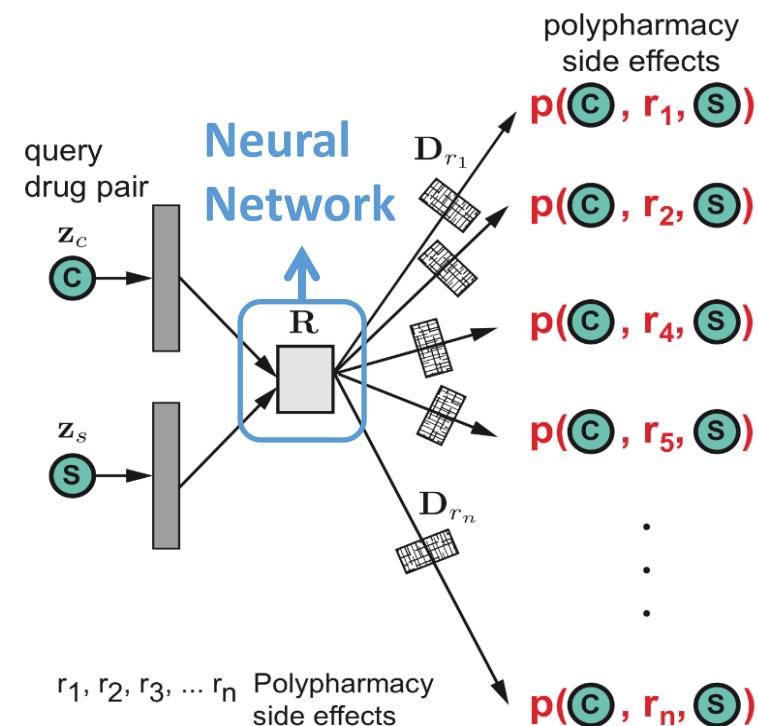
multi-class classification

- **Inference:** Use pair of computed node embeddings to make edge predictions

- **Input:** Node embeddings of query drug pairs
- **Output:** predicted edges

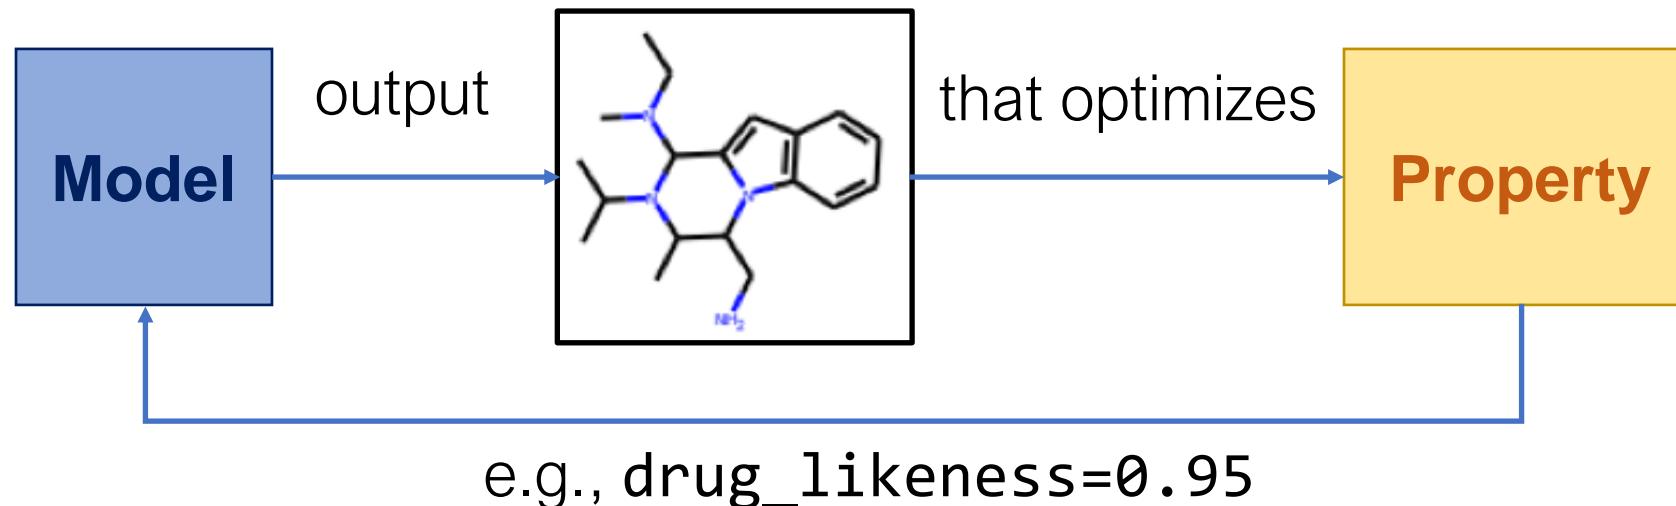


Predict possible edges with NN



# Application: Drug Discovery

**Question:** Can we learn a model that can generate **valid** and **realistic** molecules with **optimized** property scores?



Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

# Goal-Directed Graph Generation

**The goal is to generate graphs that:**

- **Optimize a given objective (High scores)**
  - e.g., drug-likeness
- **Obey underlying rules (Valid)**
  - e.g., chemical validity rules
- **Are learned from examples (Realistic)**
  - Imitating a molecule graph dataset

Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

# The Hard Part:

**The goal is to generate graphs that:**

- **Optimize a given objective (High scores)**

Only available at the end of generation

- **Obey underlying rules (Valid)**

Often not differentiable

- **Are learned from examples (Realistic)**

Requires the model to learn the distribution of graph structures

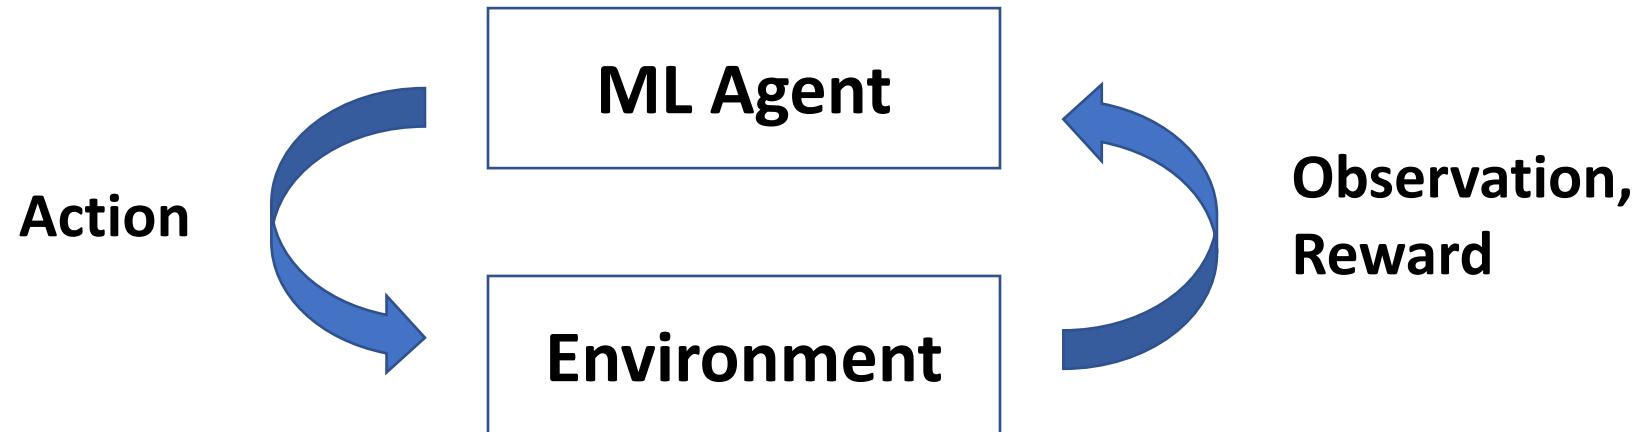
generate atom by atom  
bond by bond  
no reward in intermediate step  
only has reward in the end



Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

# Idea: Reinforcement Learning

- An ML agent **observes** the environment, takes an **action** to interact with the environment, and receives positive or negative **reward**
- The agent then **learns from this loop**
- Key idea: Agent can directly learn from environment, which is a **blackbox** to the agent



# Solution: GCPN

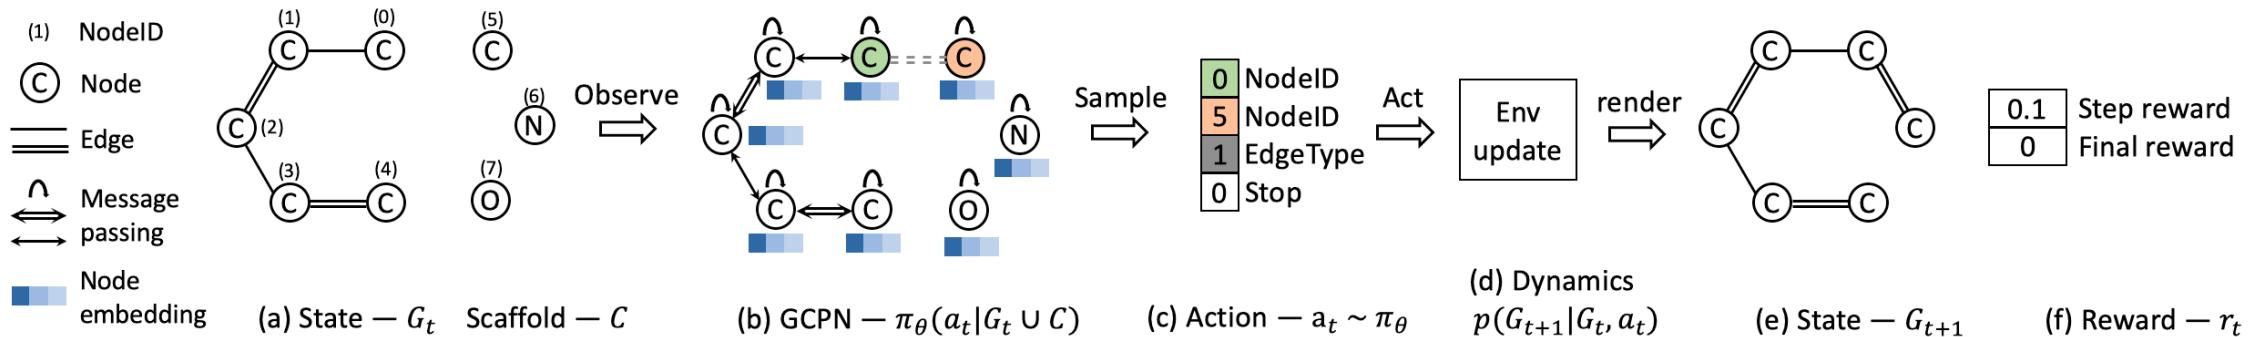
**Graph Convolutional Policy Network (GCPN)** combines **graph representation** and **reinforcement learning**

## ③ Key component of GCPN:

- ① • **Graph Neural Network** captures graph structural information
- ② • **Reinforcement learning** guides the generation towards the desired objectives
  - Generation step-by-step: every step we generate additional edges (bonds) and nodes (atoms) to attach to the generated molecule
  - Delayed reward: we only obtain the final metric (drug-likeness) at the end of the generation process
- ③ • **Supervised learning** imitates examples in given datasets

Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

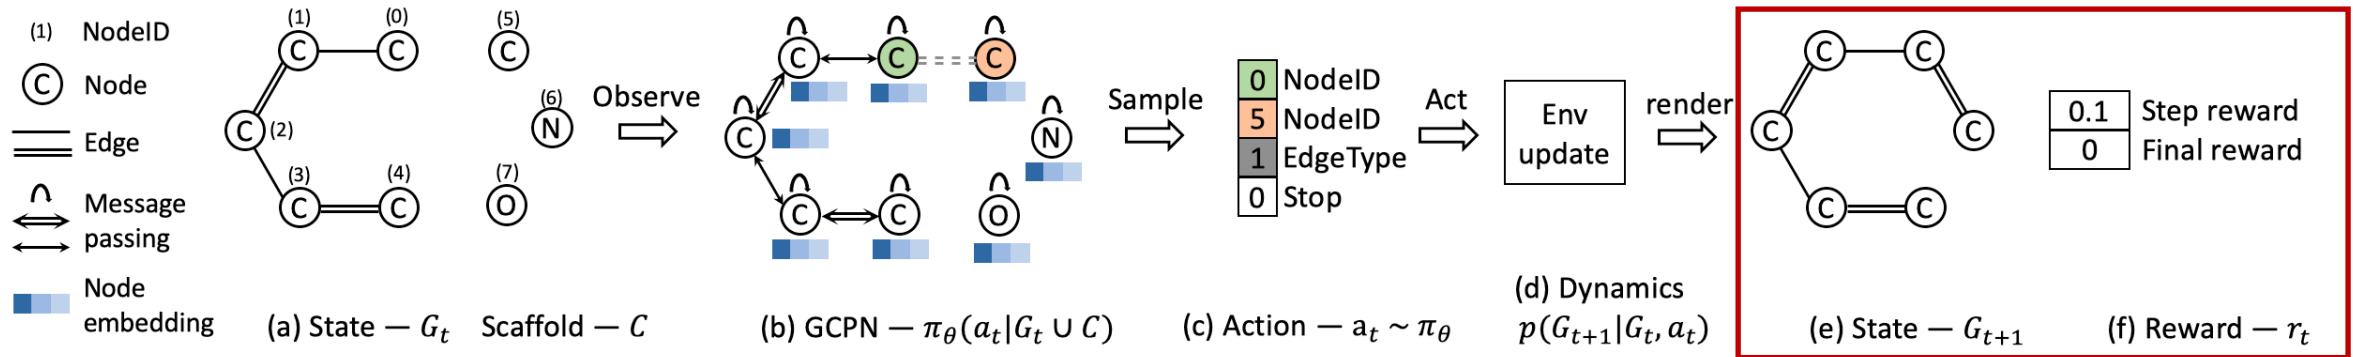
# Overview of GCPN



- **(a)** Insert nodes
- **(b,c)** Use GNN to predict which nodes to connect
- **(d)** Take action (check chemical validity)
- **(e, f)** Compute reward

Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

# How Do We Set the Reward?



- **Step reward:** Learn to take valid action
  - At each step, assign small positive reward for valid action (satisfy valency rules)
- **Final reward:** Optimize desired properties
  - At the end, assign positive reward for high desired property

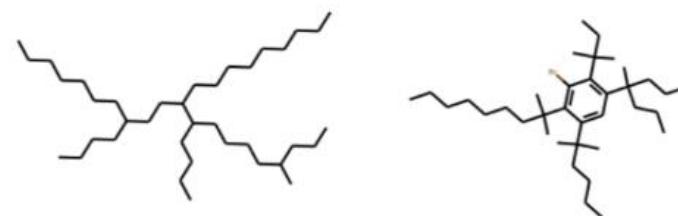
$$\text{Reward} = \text{Final reward} + \text{Step reward}$$

Graph Convolutional Policy Network for Goal-Directed Molecular Graph Generation, NeurIPS 2018.

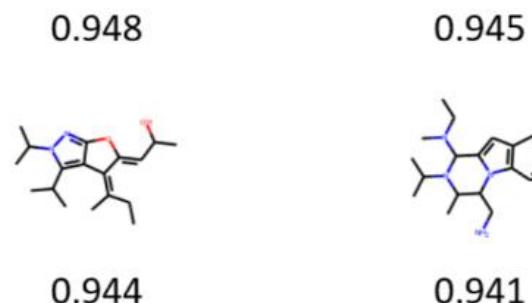
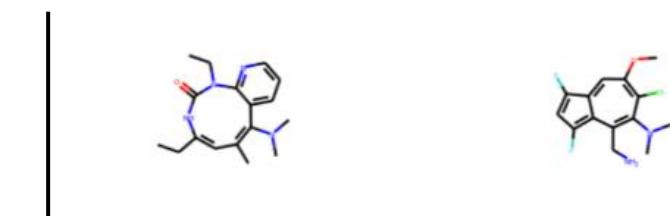
# Qualitative Results

## Visualization of GCPN graphs:

- **Property optimization** Generate molecules with high specified property score



(a) Penalized logP optimization

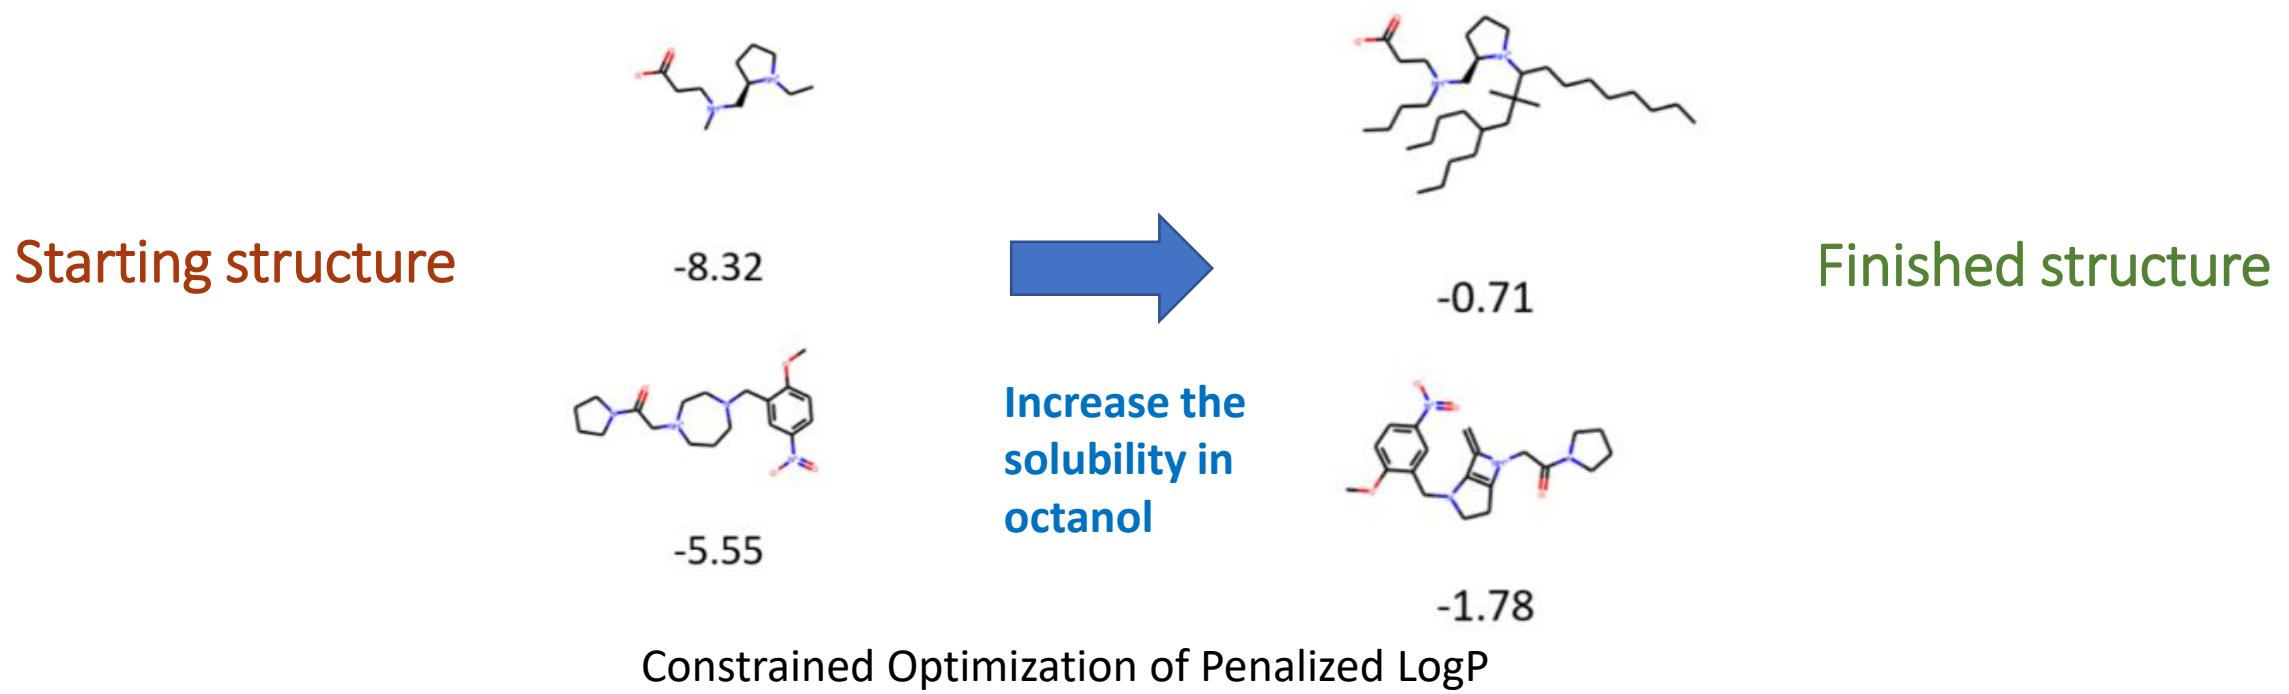


(b) QED optimization

# Qualitative Results

## Visualization of GCPN graphs:

- **Constrained optimization:** Edit a given molecule for a few steps to achieve higher property score

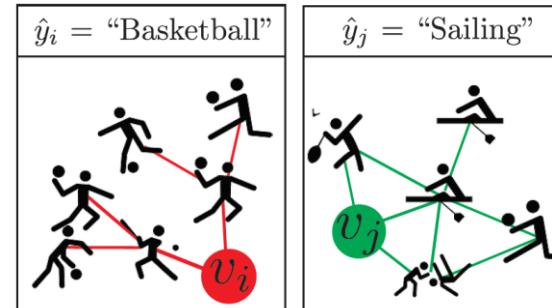


# Summary

- Many biological applications can be formulated as learning on **heterogeneous** graphs
- An ML algorithm to generate graphs (molecules) is the central problem of drug discovery
  - **Imitating** a set of given graphs
  - **Optimizing** graphs towards given goals
  - Other interesting tasks:
    - Retro-synthesis; molecule conformation; drug / molecule property prediction; prediction of the effects of gene knockouts etc.
- Many other interesting tasks!
  - Protein folding, single-cell analysis ...

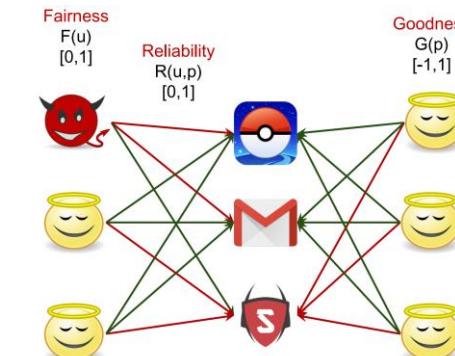
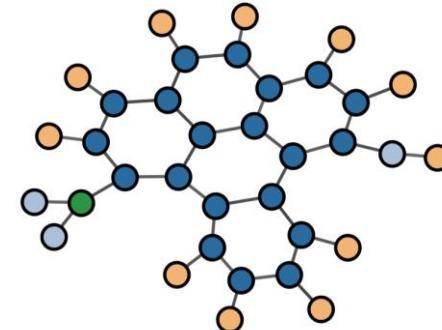
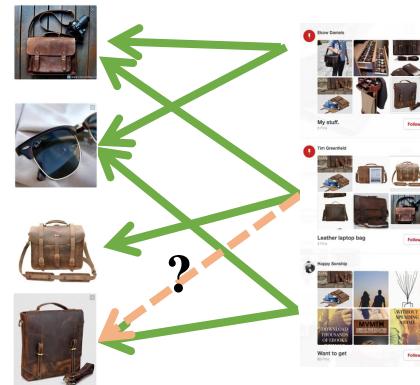
# GNN Applications

- Social Networks
- Natural Science
- Medicine
- Explainability of GNNs



# Scalable GNN Explanations

- Many questions after training GNNs:
  - Why is an item recommended to a user? Explain link prediction
  - Why is the molecule mutagenic? 诱变的 Explain graph classification
  - Why is the user classified as fraudulent? 诈骗的 Explain node classification
- Being able to answer these questions are important for domain experts



# Why is it hard

- Explain predictions for multiple tasks
  - Node classification
  - Graph classification
  - Link prediction
- Model agnostic (*post-hoc*)<sup>不相关的</sup>
  - Can be applied to all models covered in class:  
GCN, GraphSAGE, GAT etc.

# Recap: GNN Framework

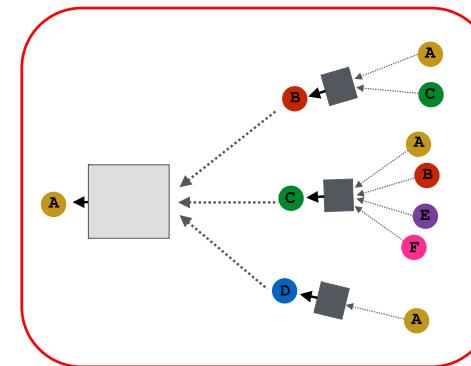
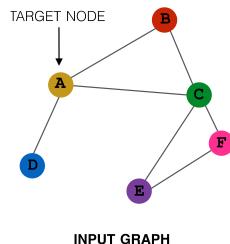
- Message Computation

$$m_{ij}^l = \text{Msg}(\mathbf{h}_i^{l-1}, \boxed{\mathbf{h}_j^{l-1}}, e_{ij})$$

Previous layer neighbor embedding

- Aggregation

$$M_i^l = \text{Agg}(\{m_{ij}^l \mid v_j \in N_{v_i}\})$$



- Representation update

$$\mathbf{h}_i^l = \text{Update}(M_i^l, \mathbf{h}_i^{l-1})$$

- Stack Multiple layers

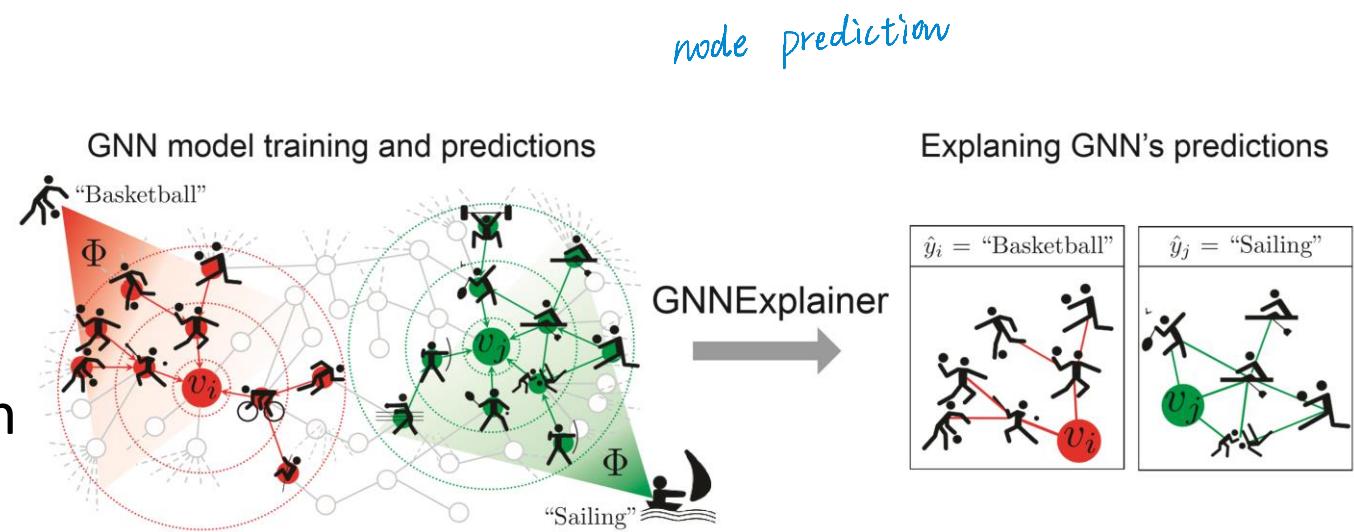
# Model Explanation

- **Training time:**

- Optimize GNN on training graphs
- Save the trained model

- **Test time:**

- Explain predictions made by the GNN
- On unseen instances (nodes, edges, graphs)

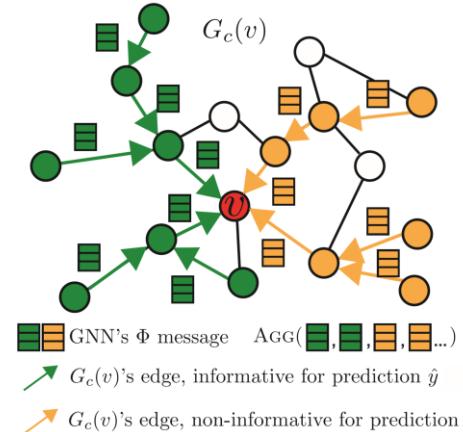


GNNExplainer: Generating Explanations for Graph Neural Networks, NeurIPS 2019

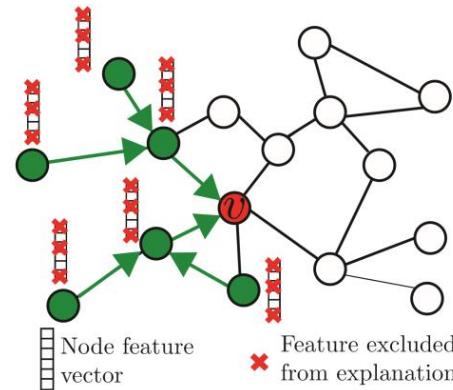
# GNNExplainer

- Message passing structure
- The importance of node features
- Explain both aspects **simultaneously**
- To explain a given node, learn
  - Important edges in its neighborhood
  - Important node feature dimensions
- **Mutual information objective**

{ structure  
feature



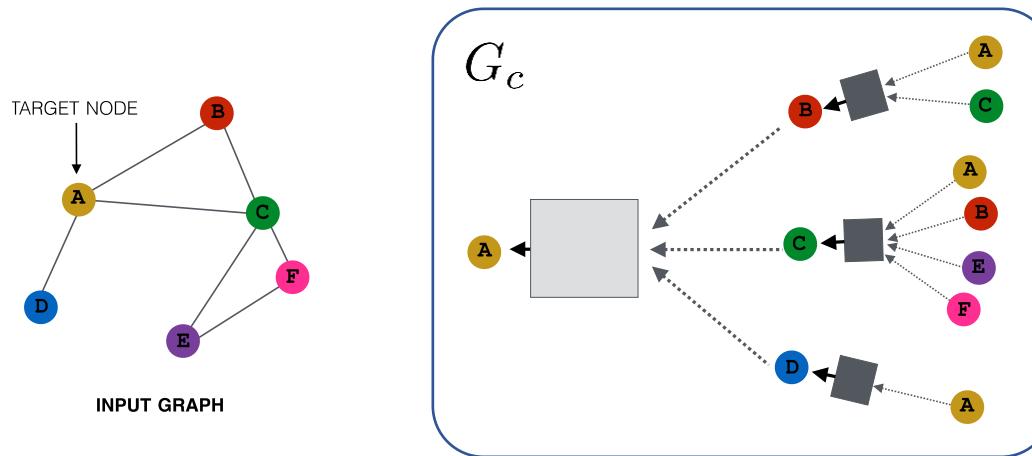
**Structural explanation**



**Feature explanation**

# GNN-Explainer Input

- Consider node classification task:



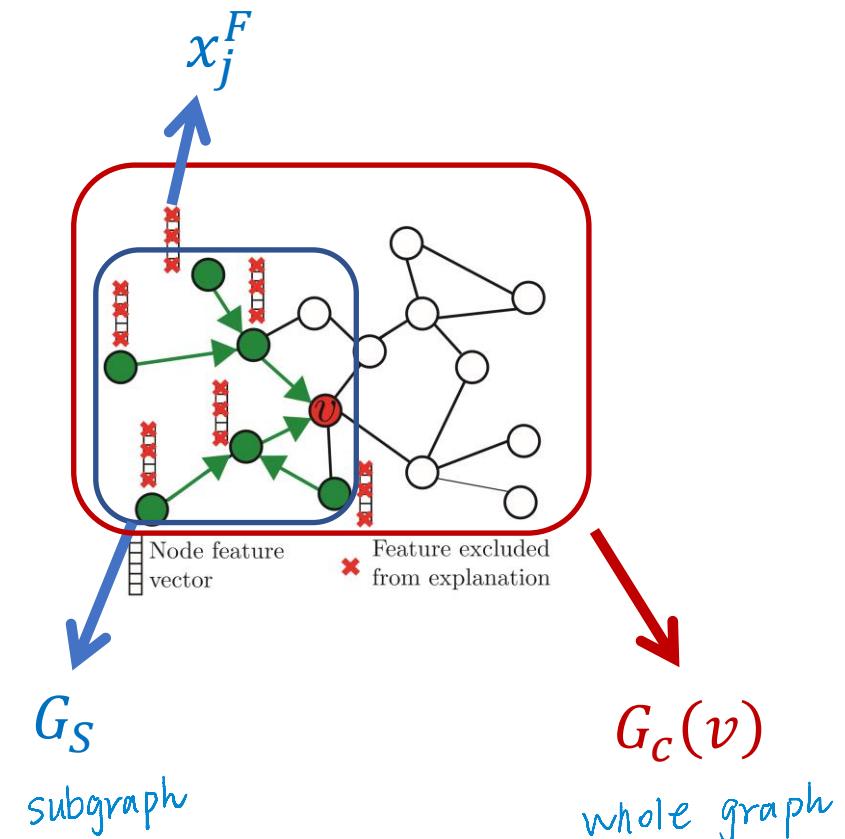
Suppose GNN predicts  
label  $\hat{y}$  for node  $v$

- Input computation graph:  $G_c(v)$
- Adjacency matrix:  $A_c(v) \in \{0,1\}^{n \times n}$
- Node Feature:  $X_c(v) = \{x_j | v_j \in G_c(v)\}$

# GNN-Explainer Output

subgraph and subgraph node features

- GNN model  $\phi$  learns  $P_\phi(Y | G_c(v), X_c(v))$
- $Y$  denotes predicted label of  $v$
- **GNNExplainer** outputs  $(G_S, X_S^F)$
- $G_S$  is a small subgraph of  $G_c(v)$  (omit  $v$ )  
*whole graph*
- $X_S^F = \{x_j^F | v_j \in G_S\}$  are features for  $G_S$
- Mask  $F$  masks out unimportant dimensions



# Experiments

- Synthetic task: **is a node part of a motif?**

	BA-Shapes	BA-Community	Tree-Cycles	Tree-Grid
Base				
Motif				
Node Features	None	$\mathcal{N}(\mu_l, \sigma_l)$ where $l$ = community ID	None	None
Explanation content	Graph structure	Graph structure Node feature information	Graph structure	Graph structure

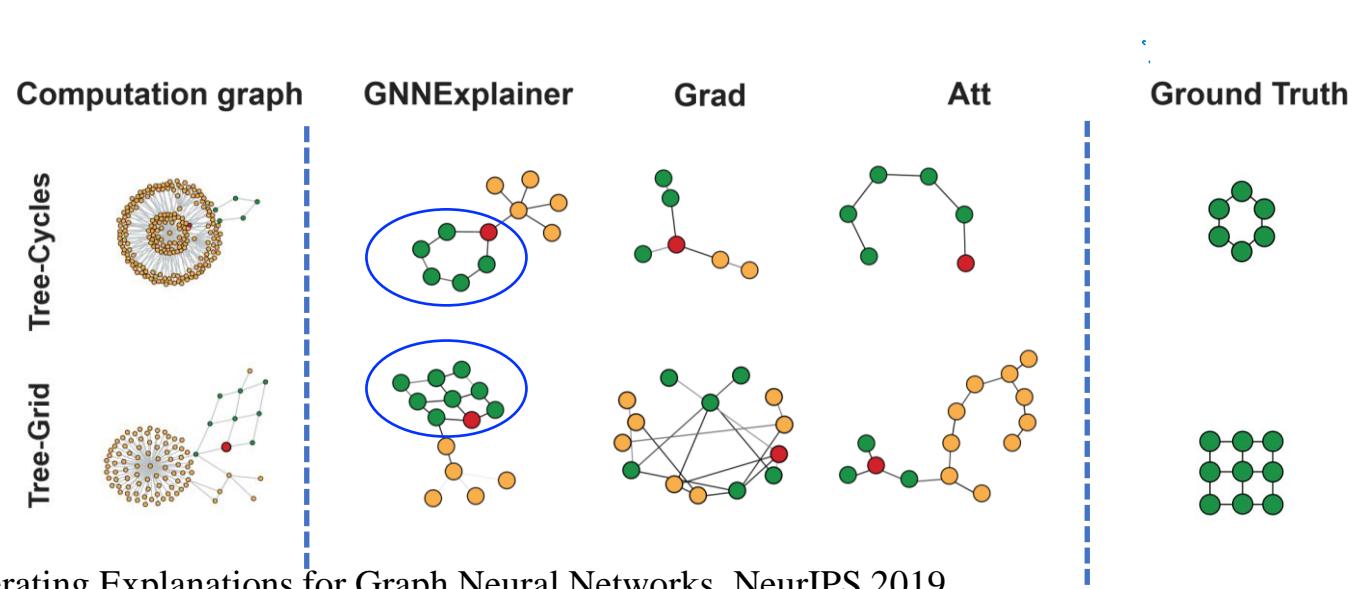
- Real-world social networks, molecules
  - **Graph classification**

# Explainability Method Comparison

- GNN saliency map based on gradients
  - Large gradient: more important
- Graph Attention Networks (GAT)
  - Edge importance indicated by attention
  - Average attention weights across layers

# GNNExplainer

- Explain node classification using GCN
- Base: **Barabasi-Albert Graph**; addon: **Cycle, Grid**



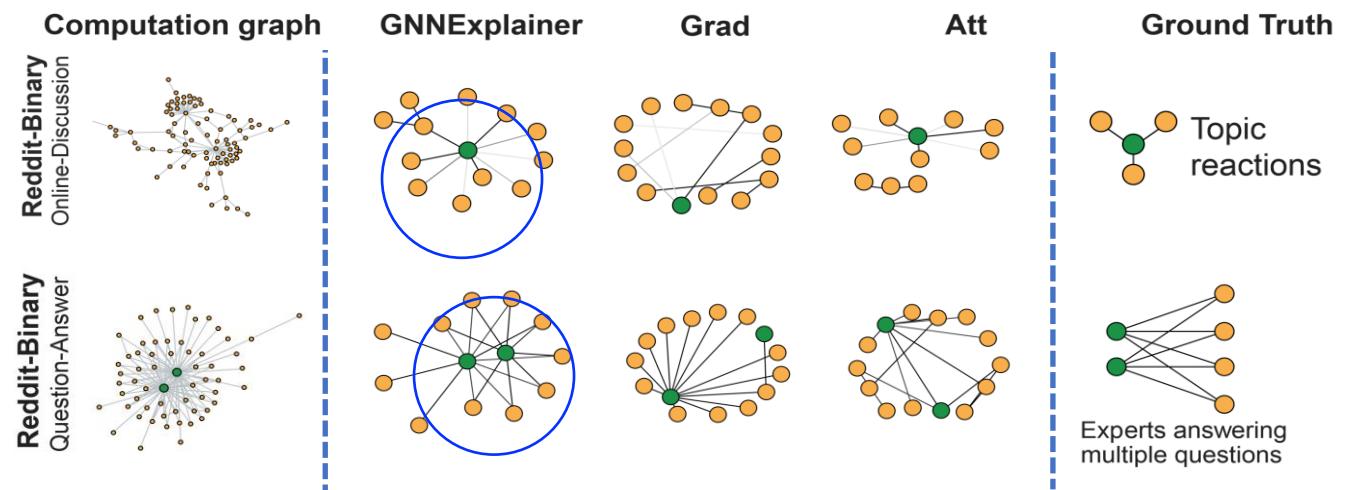
GNNExplainer: Generating Explanations for Graph Neural Networks, NeurIPS 2019

# GNNExplainer

## Graph Classification

**Reddit-Binary**: 2000 threads, 2 classes

Discussion community vs. QA community



GNNExplainer: Generating Explanations for Graph Neural Networks, NeurIPS 2019

# Conclusion

- We have covered a variety of graph learning applications, in domains including **online social networks**, **natural sciences** and **medicine**.
- Many of the problems were not always formulated as a graph problem!
- Interesting directions to explore
  - Novel ways to **model** problems through the angle of graph learning
  - **Scalable GNN algorithms** for extremely large datasets
  - **AutoML, explainability** for GNN applications

# Deep Learning on Graph-Structured Data

- New course on **Deep Learning on Graph-structured Data** (Fall 2022)
- The scope will be primarily focused on **theory, algorithms** and **applications** related to GNNs and geometric deep learning
  - We also encourage students who primary work on ML for visual, text or other forms of data to explore potential incorporation of graph techniques in their research problems